

December 6, 2020
DRAFT

Extending Abstract Effects with Bounds and Algebraic Handlers

Anlun Xu

November 2020

School of Computer Science
Carnegie Mellon University
Pittsburgh, PA 15213

Thesis Committee:

*Submitted in partial fulfillment of the requirements
for the degree of Master of Science.*

Copyright © 2020 Anlun Xu

December 6, 2020
DRAFT

Keywords: Effect Systems

December 6, 2020
DRAFT

Dedication

Abstract

Effect systems have been a subject of active research for nearly four decades, with the most notable practical example being checked exceptions in programming languages such as Java. The work on effects can be divided into two strands: The *restrictive* approach (e.g., Java’s checked exceptions), which takes effects that are already built into the language—such as reading and writing state or exceptions—and provides a way to restrict them. And the *denotational* approach, which includes algebraic effects, defines the semantics of computational effects based on primitives. While there are many existing restrictive or denotational effect system, they are rarely designed with scalability in mind. In this thesis, we design multiple effect systems around the idea of making effect system scalable when developing large and complex softwares. The first part of our work is a restrictive path-dependent effect system that provide a granular effect hierarchy by allowing abstract effect members to be bounded. This thesis presents a full formalization of the effect-system, and provides an implementation as a part of the Wyvern programming language. The second part of our work presents a denotational effect-system that supports abstract algebraic effects. We give a formalization of the system and provide proofs for type soundness and dynamic semantic correctness of abstract algebraic effects.

December 6, 2020
DRAFT

Acknowledgments

Contents

1	Introduction	1
2	Background and Motivation	5
2.1	Descriptive Effect System	5
2.2	Algebraic Effects	6
2.3	Wyvern Effect Basics	6
2.3.1	Path-dependent Effects	8
2.3.2	Effect Abstraction	8
2.3.3	Effect Aggregation	9
2.3.4	Controlling FFI Effects	9
3	Bounded Abstract Effects	11
3.1	Effect Bounds	11
3.1.1	Controlling Access to UI Objects	12
3.1.2	Controlling Mutable States Using Abstract Regions	13
3.2	Formalization	14
3.2.1	Object-Oriented Core Syntax	14
3.2.2	Modules-to-Objects Translation	15
3.2.3	Well-formedness	16
3.2.4	Static Semantics	16
3.2.5	Subtyping	19
3.2.6	Dynamic Semantics and Type Soundness	22
4	Abstract Algebraic Effects Via Embedding	25
4.1	Background and Motivation	25
4.2	Core Calculus	26
4.2.1	Syntax	26
4.2.2	Agent Specific Type Information	28
4.3	Operational Semantics	29
4.4	Static Semantics	31
4.4.1	Typing Rules	31
4.4.2	Type Relations	33
4.5	Safety Properties	33
4.6	Abstraction Problem I	41

5	Algebraic Effects with Module Systems	45
5.1	Programming with Abstract Effects	45
5.2	Future Work	47
5.2.1	Parametric Polymorphism	47
5.3	Existential Type for Algebraic Effect	48
5.3.1	Syntax	48
5.3.2	Dynamic Semantics	49
5.3.3	Static Semantics	50
5.3.4	Type Relation	50
5.3.5	Type Safety	50
5.4	Wyvern with Algebraic Effects and Handlers	51
5.4.1	Syntax	51
5.4.2	Operational Semantics	52
5.4.3	Typing Rules	52
5.4.4	Type Safety	52
6	Related Work and Conclusion	55
6.1	Related Work	55
6.1.1	Restrictive Effect Systems	55
6.1.2	Bounded Effect Polymorphism.	55
6.1.3	Subeffecting.	55
6.1.4	Path-dependent Effects	55
6.1.5	Abstract Algebraic Effects	56
	Bibliography	59

List of Figures

2.1	A type and a module implementing the logging facility in the text-editor application.	7
2.2	The type of the file resource.	7
3.1	Wyvern's object-oriented core syntax.	14
3.2	A simplified translation of the <code>logger</code> module from Fig. 2.1 into Wyvern's object-oriented core.	15
3.3	Wyvern well-formedness rules.	17
3.4	Wyvern static semantics.	18
3.5	Wyvern subeffecting rules.	19
3.6	Rules for determining the size of effect definitions.	20
3.7	Wyvern subtyping rules.	21
3.8	Algorithmic Subtyping	21
3.9	Wyvern's object-oriented core syntax with dynamic forms.	22
3.10	Wyvern static semantics affected by dynamic semantics.	23
3.11	Wyvern dynamic semantics.	24
4.1	Syntax for multi-agent calculus	27
4.2	Operational Semantics for Expressions	29
4.3	Operational Semantics for Computations	30
4.4	Static Semantics	32
4.5	Type Relations	33
4.6	Definition of $\approx_{op,op'}$	39
5.1	Syntax for Existential Effects	49
5.2	Additional Dynamic Semantics for Existential Type	49
5.3	Additional Static Semantics for Existential Effects	50
5.4	Wyvern's object-oriented core syntax with algebraic effects and handlers	52
5.5	Dynamic semantics for Wyvern extended with algebraic effects	53
5.6	Typing rules for Wyvern extended with algebraic effects	53
5.7	Subtyping rules for Wyvern extended with algebraic effects	54

List of Tables

Chapter 1

Introduction

Effect systems have been a subject of active research for nearly four decades, with the most notable practical example being checked exceptions in programming languages such as Java. According to Filinski [7], there are two different views on modeling computational effects in programs: The denotational approach and the restrictive approach.

The denotational approach describes how effectful programs can be translated into a pure program, which can then be evaluated using the standard semantics for pure programs. Works by Moggi [22] shows that features from imperative computations, such as exceptions or mutable states, can be mimicked by monads in a pure program. Algebraic effects and handlers [26] are the latest development in this strand of work. Algebraic effects and handlers can express a wide range of computation effects such as nondeterminism, concurrency, state, input/output [26]. Comparing to the traditional approach that uses general monads, algebraic effects have the advantage of being freely composable. Therefore, algebraic effects are recently gaining popularity as an approach to model effects in the purely functional setting.

Alternatively, in a restrictive setting of computation effects, the computation effects are considered to be built into the language. Rather than building up new behaviors, the effect systems aim to classify and restrict the use of existing effectful behavior in a language, such as reads and writes to memory, and checked exceptions. Restrictive effect systems are widely used for reasoning about security [30], memory effects [18], and concurrency [3, 5, 6].

Abstraction: A requirement for scalable effect systems

Unfortunately, effect systems have not been widely adopted, other than checked exceptions in Java, a feature that is widely viewed as problematic [31]. The root of the problem is that existing effect systems do not provide adequate support for scaling to programs that are larger and have complex structure. Any adequate solution must support *effect abstraction* and *effect composition*.

Abstraction is key to achieving scale in general, and a principal form of abstraction is abstract types [20]. There are many existing works that achieve information hiding using abstract types, such as SML signatures and abstract type members in Scala [23]. Typically, a module system allows each module to choose what names and entities to export, and what to keep hidden. The exported interface typically does not reveal the detail of the implementation of a module. By hiding the implementation details, the programmer of the module can be certain that the invariants within the module cannot be broken by the client. Analogously to type abstraction,

we define *effect abstraction* as the ability to define higher-level effects in terms of lower-level effects, and potentially to *hide* that definition from clients of an abstraction.

In large-scale systems, abstraction should be *composable*. For example, a database component might abstract `file.Read` further, exposing it as a higher-level `db.Query` effect to clients. Clients of the database should be oblivious to whether `db.Query` is implemented in terms of a `file.Read` effect or a `network.Access` effect (in the case that the backend is a remote database).

Design of a restrictive effect system in Wyvern

This thesis presents a novel and scalable effect-system design that supports bounded effect abstraction that extends the effect system presented by Melicher et al. The abstraction facility of our effect-system is inspired by type members in languages such as Scala. Just as Scala objects may define type members, in our effect calculus, any object may define one or more *effect members*. An effect member defines a new effect in terms of the lower-level effects that are used to implement it. The set of lower-level effects may be empty in the base case or may include low-level effects that are hard-coded in the system. Type ascription can enable information hiding by concealing the definition of an effect member from the containing object’s clients. In addition to completely concealing the definition of an effect, our calculus provides bounded abstraction, which exposes upper or lower bounds of the definition of an effect, while still hiding the definition of it.

Effect polymorphism is a form of parametric polymorphism that allows functions or types to be implemented generically for handling computations with different effects [18]. In systems at a larger scale, there are various possible effects, and each program component may cause different effects. With effect polymorphism, we can write general code that handles objects with different effects, thereby reducing the amount of replicated code. In practice, we have found that to make effects work well with modules, it is essential to extend effect polymorphism by assigning bounds to effect parameters. We therefore introduce *bounded abstract effects*, which allows programmers to define upper and lower bounds both on abstract effects and on polymorphic effect parameters.

Just as Scala’s type members can be used to encode parametric polymorphism over types, our effect members and their bounds double as a way to provide bounded effect polymorphism. We follow numerous prior Scala formalisms in including polymorphism via this encoding rather than explicitly; this keeps the formal system simpler without giving up expressive power.

Design of a denotational effect system with effect abstraction

This thesis presents a core calculus that supports algebraic effects. The calculus extends simply typed lambda calculus with algebraic effect operations and handlers and provides the ability to define abstract algebraic effects. Similar to the restrictive effect system, algebraic effects types can be defined in terms of lower-level effects. Effect abstraction in this system ensures that the client of an abstract effect type is not aware of the lower-level effects that implements the abstract effect. Consequently, the client of an abstract effect type would not be able to handle the computation that causes the abstract effect, whose operations are hidden.

Different from the restrictive effect system in Wyvern, which describes the built in effectful behavior in the language and does not affect the dynamic semantics of the program, the semantics of algebraic effects operations depend on the handler that encapsulates the operation during evaluation. Therefore, the effect system needs to ensure the abstraction does not break during the evaluation of a program. This problem was originally discovered by Biernacki et al. [2], who

solved the problem using the technique of coercions. In this paper, we propose the technique of agent-based reasoning originally designed by Grossman et al. [10] as a solution of the problem. The benefit of this approach is that by explicitly dividing modules with hidden information into agents, it is able to support syntactic proof for properties on abstract algebraic effects.

Outline and Contributions. Chapter introduces the background of both restrictive and denotational effect systems, and discusses the basics of the Wyvern effect system, after which we describe the main contributions of our paper:

- A design of a more expressive effect system for Wyvern. Specifically, ours is the first system to provide the programmer with a general form of bounded effect polymorphism and bounded effect abstraction, supporting upper and lower bounds that are other arbitrary effects. (Section 3.1);
- A precise, formal description of our effect system, and proof of its soundness. Our formal system shows how to generalize and enrich earlier work on path-dependent effects by leveraging the type theory of DOT (Section 3.2);
- A multi-agent calculus that supports abstraction for algebraic effects, and proof of its soundness; Our system enables a syntactic proof of the effect abstraction property (Section 4.2);
- A multi-agent calculus extended with existential effect types that demonstrates how multi-agent calculus could be derived from traditional techniques of type abstraction (Section 5.3);

The last chapter in the thesis discuss related work and conclude.

Chapter 2

Background and Motivation

2.1 Descriptive Effect System

Denotational vs. Descriptive Effects. Filinski [7] makes a distinction between two strands of work on effects. A *denotational* approach, which includes algebraic effects, defines the semantics of computational effects based on primitives. A *restrictive* approach (e.g., Java’s checked exceptions) takes effects that are already built into the language—such as reading and writing state or exceptions—and provides a way to restrict them.

Origins of Effect Systems. Effect Systems were originally proposed by Lucassen [17] to track reads and writes to memory, and then Lucassen and Gifford [18] extended this effect system to support polymorphism. Effects have since been used for a wide variety of purposes, including exceptions in Java [11] and asynchronous event handling [5]. Turbak and Gifford [30] previously proposed effects as a mechanism for reasoning about security, which is the main application that we discuss.

Prior Work on Bounded Effect Polymorphism. A limited form of bounded effect polymorphism were explored by Trifonov and Shao [29], who bound effect parameters by the resources they may act on; however, the bound cannot be another arbitrary effect, as in our system. Long et al. [16] use a form of bounded effect polymorphism internally but do not expose it to users of their system.

Path-Dependent Effects JML’s data groups [14] have some superficial similarities to Wyvern’s effect members. Data groups are identifiers bound in a type that refer to a collection of fields and other data groups. They allow a form of abstract reasoning, in that clients can reason about reads and writes to the relevant state without knowing the underlying definitions. Data groups are designed specifically to capture the modification of state, and it is not obvious how to generalize them to other forms of effects.

The closest prior work on path-dependent effects, by Greenhouse and Boyland [9], allows programmers to declare regions as members of types; this supports a form of path-dependency in read and write effects on regions. Our formalism expresses path-dependent effects based on the type theory of DOT [1], which we find to be cleaner and easier to extend with the unique bounded abstraction features of our system. Amin et al.’s type members can be left abstract or refined by upper or lower bounds, and were a direct inspiration for our work on bounded abstract

effects.

Subeffecting. Some effect systems, such as Koka [12], provide a built-in set of effects with fixed sub-effecting relationships between them. Rytz et al. [27] supports more flexibility via an extensible framework for effects. Users can plug in their own domain of effects, specifying an effect lattice representing sueffecting relationships. Each plugin is monolithic. In contrast, our effect members allow new effects to be incrementally added and related to existing effects using declared subeffect bounds.

2.2 Algebraic Effects

Algebraic Effects, Generativity, and Abstraction. Algebraic effects and handlers [25, 26] are a way of implementing certain kinds of side effects such as exceptions and mutable state in an otherwise purely functional setting. As described above, algebraic effects fall into the “denotational” rather than “descriptive” family of effects work; these lines of work are quite divergent, and it is often unclear how to translate technical ideas from one setting to the other. However, certain papers explore parallels to our work, despite the major contextual differences.

Bračevac et al. [5] use algebraic effects to support asynchronous, event-based reactive programs. They need to use a different algebraic effect for each join operation that correlates events; thus, they want effects to be generative. This generativity is at a per-module level, however, whereas our work supports per-object generativity.

Zhang and Myers [32] describe a design for algebraic effects that preserves abstraction in the sense of parametric functions: if a function does not statically know about an algebraic effect, that effect tunnels through that function. This is different from our form of abstraction, in which the definition of an effect is hidden from clients.

Biernacki et al. [2] discuss how to abstract algebraic effects using existentials. The setting of algebraic effects makes their work quite different from ours: their abstraction hides the “handler” of an effect, which is a dynamic mechanism that actually implements effects such as exceptions or mutable state. In contrast, our work allows a high-level effect to be defined in terms of zero or more lower-level effects, and our abstraction mechanism allows the programmer to hide the lower-level effects that constitute the higher-level effect. Our system, unlike algebraic effect systems, is purely static. We do not attempt to implement effects, but rather give the programmer a system for reasoning about side effects on system resources and program objects. It is not clear that defining a high-level effect that encapsulates multiple low-level events is sensible in the setting of algebraic effects, since this would require merging effect implementations that could be as diverse as mutable state and exception handling. It is also not clear how Biernacki et al.’s abstraction of algebraic effects could apply to the security scenarios we examine in Section ??, since some of our scenarios rely critically on abstracting lower-level events as higher-level ones.

2.3 Wyvern Effect Basics

Consider the code in Fig. 2.1 that shows a type and a module implementing the logging facility of the text-editor application. In the given implementation of the `Logger` type, the `logger`

```
1 resource type Logger
2   effect ReadLog
3   effect UpdateLog
4   def readLog(): {this.ReadLog} String
5   def updateLog(newEntry: String): {this.UpdateLog} Unit
6
7 module def logger(f: File): Logger
8 effect ReadLog = {f.Read}
9 effect UpdateLog = {f.Append}
10 def readLog(): {ReadLog} String = f.read()
11 def updateLog(newEntry: String): {UpdateLog} Unit = f.append(newEntry)
```

Figure 2.1: A type and a module implementing the logging facility in the text-editor application.

```
1 resource type File
2   effect Read
3   effect Write
4   effect Append
5   ...
6   def read(): {this.Read} String
7   def write(s: String): {this.Write} Unit
8   def append(s: String): {this.Append} Unit
9   ...
```

Figure 2.2: The type of the file resource.

module accesses the log file.¹ All modules of type `Logger` must have two methods: the `readLog` method that returns the content of the log file and the `updateLog` method that appends new entries to the log file. In addition, the `Logger` type declares two *abstract* effects, `ReadLog` and `UpdateLog`, that are produced by the corresponding methods. These effects are abstract because they are not given a definition in the `Logger` type, and so it is up to the module implementing the `Logger` type to define what they mean. The effect names are user-defined, allowing the choice of meaningful names.

The `logger` module implements the `Logger` type. To access the file system, an object of type `File` (shown in Fig. 2.2) is passed into `logger` as a parameter. The `logger` module's effect declarations are those of the `Logger` type, except now they are *concrete*, i.e., they have specific definitions. The `ReadLog` effect of the `logger` module is defined to be the `Read` effect of the `File` object, and accordingly, the `readLog` method, which produces the `ReadLog` effect, calls `f`'s `read` method. Similarly, the `UpdateLog` effect of the `logger` module is defined to be `f.Append`, and accordingly, the `updateLog` method, which produces the `UpdateLog` effect, calls `f`'s `append` method. In general, effects in a module or object definition must always be concrete, whereas effects in a type definition may be either abstract or concrete.

¹The keyword `resource` in the type definition indicates that the implementations of this type may have state and may access system resources; this is orthogonal to effect checking.

2.3.1 Path-dependent Effects

Effects are members of objects², so we refer to them with the form `variable.EffectName`, where `variable` is a reference to the object defining the effect and `EffectName` is the name of the effect. For example, in the definition of the `ReadLog` effect of the `logger` module, `f` is the variable referring to a specific file and `Read` is the effect that the `read` method of `f` produces. This conveniently ties together the resource and the effects produced on it (which represent the operations performed on it), helping a software architect or a security analyst to reason about how resources are used by any particular module and its methods. For example, when analyzing the effects produced by `logger`’s `readLog` method, a security analyst can quickly deduce that calling that method affects the file resource and, specifically, the file is read, simply by looking at the `Logger` type and `logger`’s effect definitions but not at the method’s code. Furthermore, these properties can be automatically checked with an idiom of use: In addition to directly looking at the effect annotation of the method of the `logger` module, the security analyst may write client code that specify the effect that the `logger` module is allowed to have. If the `logger` module accesses system resources outside of the specified effect set, then the compiler would automatically reject the program.

Because an effect includes a reference to an object instance, our effect system can distinguish reads and writes on different file instances. If the developer does not want this level of precision, it is still possible to declare effects at the module level (i.e., as members of a `fileSystem` module object instance), and to share the same `Read` and `Write` effects (for example) across all files in `fileSystem`.

The basic mechanisms of path-dependence are borrowed from Scala and have been shown to scale well in practice. These mechanisms come from the Dependent Object Types (DOT) calculus [1], a type theory of Scala and related languages (including Wyvern). In our system, effects, instead of types are declared as members of objects.

2.3.2 Effect Abstraction

An important and novel feature of our effect system design is the support for *effect abstraction*. Effect abstraction is the ability to define higher-level effects in terms of lower-level effects and potentially to hide that definition from clients of an abstraction. In the logging example above, through the use of abstraction, we “lifted” low-level resources such as the file system (i.e., the `Read` and `Append` effects of the file) into higher-level resources such as a logging facility (i.e., the `ReadLog` and `UpdateLog` effect of the logger) and enabled application code to reason in terms of effects on those higher-level resources when appropriate.

Effect abstraction has several concrete benefits. First, it can be used to distinguish different uses of a low-level effect. For example, `system.FFI` describes any access to system resources via calls through the foreign function interface (FFI), but modules that define file and network I/O can represent these calls as different effects, which enables higher-level modules to reason about file and network access separately. Second, multiple low-level effects can be aggregated into a single high-level effect to reduce effect specification overhead. For instance, the `db.Query` effect might include both `file.Read` and `network.Access` effects. Third, by keeping an effect

²Modules are an important special case of objects

abstract, we can hide its implementation from clients, which facilitates software evolution: code defining a high-level effect in terms of lower-level ones can be rewritten (or replaced) to use a different set of lower-level effects without affecting clients (more on this in Section ??).

2.3.3 Effect Aggregation

Wyvern’s effect-system design allows reducing the effect-annotation overhead by aggregating several effects into one. For example, if, to update the log file, the `logger` module needed to first read the file and then write it back, the `UpdateLog` effect would consist of two effects: a file read and a file write. In other effect systems, this change may make effects more verbose since all the methods that call the `updateLog` method would need to be annotated with the two effects. However, effect aggregation allows us to define the `UpdateLog` effect to be the two effects and then use `UpdateLog` to annotate the `updateLog` method and all methods that call it:

```
module def logger(f: File): Logger
  effect UpdateLog = {f.Read, f.Write}
  def updateLog(newEntry: String): {this.UpdateLog} Unit
  ...
```

This way we need to use only one effect, `UpdateLog`, instead of two, in method effect annotations, thus reducing the effect-annotation overhead. Because more code may add more effects, larger software systems might experience a snowballing of effects, when method annotations have numerous effects in them.

2.3.4 Controlling FFI Effects

Wyvern programs access system resources via calls to other programming languages, such as Java and Python, i.e., through a foreign function interface (FFI). To monitor and control the effects caused by FFI calls, we enforce that all functions from other programming languages, when called within Wyvern, are annotated with the `system.FFI` effect.

As was mentioned in Section 2.3.2, the `system.FFI` effect is an effect that describes function calls through an FFI. Since every function call through FFI has this effect, the access to system resources via FFI is guaranteed to be monitored. `system.FFI` is the lowest-level effect in the effect system which can be used to build other higher-level effects. The programmer can lift `system.FFI` to higher-level effects and reason about those higher-level effects instead.

For example, Wyvern’s import mechanism works by loading an object in a static field of a Java class, and the following code imports a field of a Java class that helps to implement file IO:

```
import java:wyvern.stdlib.support.FileIO.file
```

The file object is itself of type `FileIO`. And `FileIO` has this method, among others:

```
public void writeStringIntoFile(String content, String filename) throws
  IOException { ... }
```

In Wyvern, there is a type `wyvern.stdlib.support.FileIO` as well as an object `file` (of that type) that gets added to the scope as a result of the import above. The type has the following member, corresponding to the method above:

```
def writeStringIntoFile(content:String, filename:String): { system.FFI }  
Unit
```

Here, the `system.FFI` effect was added to the signature because this is a function that was imported via the FFI. The Wyvern file library that uses the `writeStringIntoFile` function abstracts this `system.FFI` effect into a library-specific `FileIO.Write` effect.

Chapter 3

Bounded Abstract Effects

3.1 Effect Bounds

Our effect system also gives the programmer the ability to define a subtyping hierarchy of effects via effect bounds. To define the hierarchy, the programmer gives the effect member an upper bound or a lower bound, hiding the definition of the effect from the client.

For example, consider the type `BoundedLogger` which has the same method declarations and effect members as the type `Logger` in Fig. 2.1, except the `ReadLog` and `UpdateLog` effects are upper-bounded by the corresponding effects in the `fileSystem` module:

```
resource type BoundedLogger
  effect ReadLog <= {fileSystem.Read}
  effect UpdateLog <= {fileSystem.Append}
  ... // same as in the type Logger in Fig. 2
```

Any object implementing type `BoundedLogger` may have an effect member `ReadLog` which is *at most* `fileSystem.Read`. This allows programmers to compare the `ReadLog` effect with other effects, while keeping its definition abstract. For instance, a library can provide two implementations of `BoundedLogger`, including an effectless logger in which the effects `ReadLog` and `UpdateLog` are empty sets, and an effectful logger in which `ReadLog` and `UpdateLog` are defined as effects in the `fileSystem` module. The library's clients then can annotate the effects of both implementations with `fileSystem.Read` and `fileSystem.Append` according to the effect hierarchy, without the need to know the exact implementation of the two instances.

Effect hierarchy can also be constructed using lower bounds. For example, consider the following type for I/O modules that supports writes:

```
type IO
  effect Write >= {system.FFI}
  def write(s: String): {this.Write} Unit
```

Since I/O is done using the foreign function interface (FFI), the `Write` effect is *at least* the `system.FFI` effect. Similar to providing an upper bounded on effects, this type does not specify the exact definition of the `Write` effect, and implementations of this type can define `Write` as an effect set with more effects than `{system.FFI}`.

The effect hierarchy achieved by bounding effect members is supported by the subtyping

relations of our effect system (Sections 3.2.5 and 3.2.5). If a type has an effect member with more strict bounds than another type, then the former type is a subtype of the latter type. For example, when a logger with the effect member `Read <= {fileSystem.Read}` is expected, we can pass in a logger with `Read = {}` because the definition as an empty set is more strict than an upper bound.

The following two case studies demonstrates the expressiveness of the effect hierarchy:

3.1.1 Controlling Access to UI Objects

This main idea of the work of [8] is to control the access of user interface (UI) framework methods so that unsafe UI methods can only be called on the UI thread. There are three different method annotations `@SafeEffect`, `@UIEffect`, and `@PolyUIEffect`, where

1. `@SafeEffect` annotates methods that are safe to run on any thread,
2. `@UIEffect` annotates methods that is only callable on UI thread, and
3. `@PolyUIEffect` annotates methods whose effect is polymorphic over the receiver type's effect parameter.

In Wyvern, we can model `@UIEffect` as a member of the UI module, for example:

```
type UILibrary
  effect UIEffect >= {system.FFI}
  def unsafeUIMethod1(): {this.UIEffect} Unit
  def unsafeUIMethod2(): {this.UIEffect} Unit
  ...
```

This way, any client code of an UI library that calls UI methods will have the `uilibrary.UIEffect` effect.

An interface could be used for UI-effectful or UI-safe work. To accommodate such flexibility, *Java_{UI}* introduced the `@PolyUIType` annotation. For example, a `Runnable` interface which can be UI-safe or UI-unsafe is declared as

```
@PolyUIType public interface Runnable {
  @PolyUIEffect void Run();
}
```

Whether the method `Run()` will have a UI effect depends on an annotation when the type is instantiated. For example:

```
@Safe Runnable s = ....;
s.run(); // is UI safe
@UI Runnable s = ....;
s.run(); // has UI effect
```

In Wyvern, such polymorphic interface can be created by defining the interface with a bounded effect member:

```
type Runnable
  effect Run <= {uiLibrary.UIEffect}
  def run(): {this.Run} Unit
```

This type ensures that the `run` method is safe to be called on the UI thread. Moreover, if an instance of `Runnable` does not have `UIEffect`, it can be ascribed with the type `SafeRunnable`, which is a subtype of `Runnable`:

```
type SafeRunnable
  effect Run = {}
  def run(): {this.Run} Unit
```

This indicates that `run` is safe to be called on any thread.

3.1.2 Controlling Mutable States Using Abstract Regions

[9] proposed a region-based effect system which describes how state may be accessed during the execution of some program component in object-oriented programming languages. One example of the usage of regions is as follows:

```
class Point {
  public region Position;
  private int x in Position;
  private int y in Position;
  public scale(int sc) reads nothing writes Position {
    x *= sc;
    y *= sc;
  }
}
```

The two variables `x` and `y` are declared inside a region `Position`. For each region, there can be two possible effects: read and write. The `scale` method has the effect of writing on the region `this.Position`.

To achieve access control on regions in Wyvern, we need to keep track of the read and write effect on each variable in a region. We declare the resource type `Var` representing a variable wrapper.

```
resource type Var[T]
  effect Read
  effect Write
  def set (x: T): {this.Write} Unit
  def get (): {this.Read} T
```

Since the `set` and `get` methods are annotated with the corresponding effects and there is no exposed access to the variable that holds the value, the two methods protect the access to the variable inside the type `Var`. To avoid code boilerplate, this wrapper type can be added as a language extension. The `Point` example above can be rewritten in Wyvern as:

```
resource type Point
  val x: Var[Int]
  val y: Var[Int]
  effect Read >= {this.x.Read, this.y.Read}
  effect Write >= {this.x.Write, this.y.Write}
  def scale(sc: Int): {this.Write} Unit
```

We can also extend the type `Point` to `3DPoint` in the following way:

```
resource type 3DPoint
  val x: Var[Int]
  val y: Var[Int]
  val z: Var[Int]
  effect Read = {this.x.Read, this.y.Read, this.z.Read}
  effect Write = {this.x.Write, this.y.Write, this.z.Write}
  def scale(sc: Int): {this.Write} Unit
```

Since the effect `Read` and `Write` in the type `Point` is declared with a lower bound, the type `3DPoint` is a subtype of `Point`.

3.2 Formalization

As was mentioned earlier, Wyvern modules are first class and are, in fact, objects since they are only syntactic sugar on top of Wyvern’s object-oriented core and can be translated into objects. The translation has been described in detail previously [19], and here we provide only some intuition behind it. In this section, we start with describing the syntax of Wyvern’s object-oriented core, then present an example of the module-to-object translation, followed by a description of Wyvern’s static semantics and subtyping rules. Furthermore, we present the dynamic semantics and the type soundness theorems. Last but not least, we provide the definitions on authority and discuss why they are useful for security analysis on programs written in Wyvern.

3.2.1 Object-Oriented Core Syntax

$e ::= x$	$d ::= \text{def } m(x : \tau) : \{\varepsilon\} \tau = e$	$\sigma ::= \text{def } m(x : \tau) : \{\varepsilon\} \tau$
$ \text{new } (x \Rightarrow \bar{d})$	$ \text{var } f : \tau = x$	$ \text{var } f : \tau$
$ e.m(e)$	$ \text{effect } g = \{\varepsilon\}$	$ \text{effect } g$
$ e.f$	$\varepsilon ::= \bar{x}.g$	$ \text{effect } g \geq \{\varepsilon\}$
$ e.f = e$	$\tau ::= \{x \Rightarrow \bar{\sigma}\}$	$ \text{effect } g \leq \{\varepsilon\}$
	$\Gamma ::= \emptyset \mid \Gamma, x : \tau$	$ \text{effect } g = \{\varepsilon\}$

Figure 3.1: Wyvern’s object-oriented core syntax.

Fig. 3.1 shows the syntax of Wyvern’s object-oriented core. Wyvern expressions include variables and the four basic object-oriented expressions: the `new` statement, a method call, a field access, and a field assignment. Objects are created by `new` statements that contain a variable x representing the current object along with a list of declarations. In our implementation, x defaults to `this` when no name is specified by the programmer. Declarations come in three kinds: a method declaration, a field, and an effect member. Method declarations are annotated with a set of effects. Object fields may only be initialized using variables, a restriction which simplifies our core language by ensuring that object initialization never has an effect. Although at first this may seem to be limiting, in fact, we do not limit the source language in this way. Side-effecting member initializations in the source language are translated to the core by wrapping the new object with a `let` expression (a discussion of which is upcoming) that defines the variable to be used in the field initialization. For example, this code:

```
new
  var x: String = f.read()
```

can be internally rewritten as:

```
let y = f.read()
in new
  var x: String = y
```

Effects in method annotations and effect-member definitions are surrounded by curly braces to visually indicate that they are sets, and each effect in an effect set is defined to be a variable representing the object on which an effect is produced, followed by a dot and the effect name. Abstract effects may be defined with an upper bound or a lower bound.

Object types are a collection of declaration types, which include method signatures, field-declaration types, and the types of effect-member declarations and definitions. Similar to the difference between the modules and their types, effects in an object must always be defined (i.e., always be concrete), whereas effects in object types may or may not have definitions (i.e., be either abstract or concrete), and may have an upper or lower bound.

3.2.2 Modules-to-Objects Translation

```
1 let logger = new(x =>
2   def apply(f : File) : {} Logger
3     new(_ =>
4       effect ReadLog = {f.Read}
5       effect UpdateLog = {f.Append}
6       def readLog() : {ReadLog} String = f.read()
7       def updateLog(newEntry : String) : {UpdateLog} Unit = f.append(newEntry)))
8 in ...// calls logger.apply(...)
```

Figure 3.2: A simplified translation of the `logger` module from Fig. 2.1 into Wyvern’s object-oriented core.

Fig. 3.2 presents a simplified translation of the `logger` module from Fig. 2.1 into Wyvern’s object-oriented core (for a full description of the translation mechanism, refer to [19]). For our purposes, the functor becomes a regular method, called `apply`, that has the return type `Logger` and the same parameters as the module functor. The method’s body is a new object containing all the module declarations. The `apply` method is the only method of an outer object that is assigned to a variable whose name is the module’s name. Later on in the code, when the `logger` module needs to be instantiated, the `apply` method is called with appropriate arguments passed in.

To aid this translation mechanism, we use the two relatively standard encodings:

$$\text{let } x = e \text{ in } e' \equiv \text{new}(_ \Rightarrow \text{def } f(x : \tau) : \tau' = e').f(e)$$

$$\text{def } m(\bar{x} : \bar{\tau}) : \tau = e \equiv \text{def } m(x : (\tau_1 \times \tau_2 \times \dots \times \tau_n)) : \tau = [x.n/x_n]e$$

The `let` expression is encoded as a method call on an object that contains that method with the `let` variable being the method’s parameter and the method body being the `let`’s body. The

multiparameter version of the method definition is encoded using indexing into the method parameters.

3.2.3 Well-formedness

Since Wyvern’s effects are defined in terms of variables, before we type check expressions, we must make sure that effects and types are well formed. Wyvern well-formedness rules are mostly straightforward and are shown in Fig. 3.3. The three judgements read that, in the variable typing context Γ , the type τ , the declaration type σ , and the effect set ε are well formed, respectively.

An object type is well formed if all of its declaration types are well formed. A method-declaration type is well formed if the type of its parameter, its return type, and the effects in its effect annotation are well formed. A field-declaration type is well formed if its type is well formed. Since an effect-declaration type has no right-hand side, it is trivially well formed. An effect-definition type is well formed if the effect set in its right-hand side is well formed. Finally, a bounded effect declaration is well formed if the upper bound or lower bound on the right-hand side is well formed. An effect set is well formed if, for every effect it contains, the definition of the effect doesn’t form a cycle, the variable in the first part of the effect is well typed and the type of that variable contains either an effect-declaration or an effect-definition type, in which the effect name matches the effect name in the second part of the effect.

The $\Gamma \vdash \text{safe}(x.g, \varepsilon)$ judgment ensures that the definition of effect $x.g$ doesn’t contain a cycle. The rules Safe-1, Safe-2, and Safe-3 are identical except the declaration of the effect type. The effect set ε memorizes a set of effects that are defined by $x.g$. The rule ensures that those effects do not appear in the definition of $x.g$, therefore eliminating cycles in effect definition.

3.2.4 Static Semantics

Wyvern’s static semantics is presented in Fig. 3.4. Expression type checking includes checking the effects that an expression may have, the set of which is denoted in a pair of curly braces between the colon and the type in the type annotation. Then, for expressions, the judgement reads that, in the variable typing context Γ , the expression e is a well-typed expression with the effect set ε and the type τ .

A variable trivially has no effects. A `new` expression also has no effects because of the fact that fields may be initialized only using variables. A new object is well typed if all of its declarations are well typed.

A method call is well typed if the expression passed into the method as an argument is well typed, if the expression the method is called on is well typed, and if the expression’s type contains a matching method-declaration type. In addition, bearing the appropriate variable substitutions, the effect set annotating the method-declaration type must be well formed, and the effect set ε in the method-call type must be a union of the effect sets of both expressions involved in the method call as well as the the effect set of the method-declaration type. The expressions that are being substituted are always the terminal runtime form, i.e., the expressions have been fully evaluated before they are substituted.

An object field read is well typed if the expression on which the field is dereferenced is well typed and the expression’s type contains a matching field-declaration type. The effects of an

$$\boxed{\Gamma \vdash \tau \text{ wf}}$$

$$\frac{\forall \sigma \in \bar{\sigma}, \Gamma, x : \{x \Rightarrow \bar{\sigma}\} \vdash \sigma \text{ wf}}{\Gamma \vdash \{x \Rightarrow \bar{\sigma}\} \text{ wf}} \text{ (WF-TYPE)}$$

$$\boxed{\Gamma \vdash \sigma \text{ wf}}$$

$$\frac{\Gamma \vdash \tau_2 \text{ wf} \quad \Gamma, x : \tau_2 \vdash \tau_1 \text{ wf} \quad \Gamma, x : \tau_2 \vdash \varepsilon \text{ wf}}{\Gamma \vdash \text{def } m(x : \tau_2) : \{\varepsilon\} \tau_1 \text{ wf}} \text{ (WF-DEF)} \quad \frac{\Gamma \vdash \tau \text{ wf}}{\Gamma \vdash \text{var } f : \tau \text{ wf}} \text{ (WF-VAR)}$$

$$\frac{}{\Gamma \vdash \text{effect } g \text{ wf}} \text{ (WF-EFFECT1)} \quad \frac{\Gamma \vdash \varepsilon \text{ wf}}{\Gamma \vdash \text{effect } g = \{\varepsilon\} \text{ wf}} \text{ (WF-EFFECT2)}$$

$$\frac{\Gamma \vdash \varepsilon \text{ wf}}{\Gamma \vdash \text{effect } g \leq \{\varepsilon\} \text{ wf}} \text{ (WF-EFFECT3)} \quad \frac{\Gamma \vdash \varepsilon \text{ wf}}{\Gamma \vdash \text{effect } g \geq \{\varepsilon\} \text{ wf}} \text{ (WF-EFFECT4)}$$

$$\boxed{\Gamma \vdash \varepsilon \text{ wf}}$$

$$\frac{\forall i, j, x_i.g_j \in \varepsilon, \Gamma \vdash \text{safe}(x_i.g_j, \{\}), \Gamma \vdash x_i : \{\} \{y_i \Rightarrow \bar{\sigma}_i\}, \text{ (effect } g_j \in \bar{\sigma}_i \vee \text{effect } g_j = \{\varepsilon_j\} \in \bar{\sigma}_i \vee \text{effect } g_j \geq \{\varepsilon_j\} \in \bar{\sigma}_i \vee \text{effect } g_j \leq \{\varepsilon_j\} \in \bar{\sigma}_i)}{\Gamma \vdash \varepsilon \text{ wf}} \text{ (WF-EFFECT)}$$

$$\boxed{\Gamma \vdash \text{safe}(x.g, \varepsilon)}$$

$$\frac{\Gamma \vdash x : \{\} \{y \Rightarrow \bar{\sigma}\}, \text{effect } g = \{\varepsilon'\} \in \bar{\sigma} \quad \forall a.b \in \{x.g\} \cup \varepsilon, a.b \notin [x/y]\varepsilon' \quad \forall c.d \in [x/y]\varepsilon', \Gamma \vdash \text{safe}(c.d, \{x.g\} \cup \varepsilon)}{\Gamma \vdash \text{safe}(x.g, \varepsilon)} \text{ (SAFE-1)}$$

$$\frac{\Gamma \vdash x : \{\} \{y \Rightarrow \bar{\sigma}\}, \text{effect } g \geq \{\varepsilon'\} \in \bar{\sigma} \quad \forall a.b \in \{x.g\} \cup \varepsilon, a.b \notin [x/y]\varepsilon' \quad \forall c.d \in [x/y]\varepsilon', \Gamma \vdash \text{safe}(c.d, \{x.g\} \cup \varepsilon)}{\Gamma \vdash \text{safe}(x.g, \varepsilon)} \text{ (SAFE-2)}$$

$$\frac{\Gamma \vdash x : \{\} \{y \Rightarrow \bar{\sigma}\}, \text{effect } g \leq \{\varepsilon'\} \in \bar{\sigma} \quad \forall a.b \in \{x.g\} \cup \varepsilon, a.b \notin [x/y]\varepsilon' \quad \forall c.d \in [x/y]\varepsilon', \Gamma \vdash \text{safe}(c.d, \{x.g\} \cup \varepsilon)}{\Gamma \vdash \text{safe}(x.g, \varepsilon)} \text{ (SAFE-3)}$$

$$\frac{\Gamma \vdash x : \{\} \{y \Rightarrow \bar{\sigma}\}, \text{effect } g}{\Gamma \vdash \text{safe}(x.g, \varepsilon)} \text{ (SAFE-4)}$$

Figure 3.3: Wyvern well-formedness rules.

object field type are those of the expression on which the field dereferencing is called.

A field assignment is well typed if the expression to which the field belongs is well typed and the expression's type has an appropriate field-declaration type, and if the expression in the right-hand side of the assignment is well typed. The effect set that a field assignment produces

$$\boxed{\Gamma \vdash e : \{\varepsilon\} \tau}$$

$$\begin{array}{c}
\frac{x : \tau \in \Gamma}{\Gamma \vdash x : \{\} \tau} \text{ (T-VAR)} \quad \frac{\forall i, d_i \in \bar{d}, \sigma_i \in \bar{\sigma}, \Gamma, x : \{x \Rightarrow \bar{\sigma}\} \vdash d_i : \sigma_i}{\Gamma \vdash \mathbf{new}(x \Rightarrow \bar{d}) : \{\} \{x \Rightarrow \bar{\sigma}\}} \text{ (T-NEW)} \\
\\
\frac{\Gamma \vdash e_1 : \{\varepsilon_1\} \{x \Rightarrow \bar{\sigma}\} \quad \mathbf{def} \ m(y : \tau_2) : \{\varepsilon_3\} \tau_1 \in \bar{\sigma} \quad \Gamma \vdash [e_1/x][e_2/y]\varepsilon_3 \text{ wf} \quad \Gamma \vdash e_2 : \{\varepsilon_2\} [e_1/x]\tau_2 \quad \varepsilon = \varepsilon_1 \cup \varepsilon_2 \cup [e_1/x][e_2/y]\varepsilon_3}{\Gamma \vdash e_1.m(e_2) : \{\varepsilon\} [e_1/x][e_2/y]\tau_1} \text{ (T-METHOD)} \\
\\
\frac{\Gamma \vdash e : \{\varepsilon\} \{x \Rightarrow \bar{\sigma}\} \quad \mathbf{var} \ f : \tau \in \bar{\sigma}}{\Gamma \vdash e.f : \{\varepsilon\} [e/x]\tau} \text{ (T-FIELD)} \quad \frac{\Gamma \vdash e : \{\varepsilon_1\} \tau_1 \quad \Gamma \vdash \tau_1 <: \tau_2 \quad \Gamma \vdash \varepsilon_1 <: \varepsilon_2}{\Gamma \vdash e : \{\varepsilon_2\} \tau_2} \text{ (T-SUB)} \\
\\
\frac{\Gamma \vdash e_1 : \{\varepsilon_1\} \{x \Rightarrow \bar{\sigma}\} \quad \mathbf{var} \ f : \tau \in \bar{\sigma} \quad \Gamma \vdash e_2 : \{\varepsilon_2\} \tau \quad \varepsilon = \varepsilon_1 \cup \varepsilon_2}{\Gamma \vdash e_1.f = e_2 : \{\varepsilon\} [e_1/x]\tau} \text{ (T-ASSIGN)}
\end{array}$$

$$\boxed{\Gamma \vdash d : \sigma}$$

$$\begin{array}{c}
\frac{\Gamma, x : \tau_1 \vdash e : \{\varepsilon_2\} \tau_2 \quad \Gamma, x : \tau_1 \vdash \varepsilon_1 \text{ wf} \quad \Gamma, x : \tau_1 \vdash \varepsilon_2 <: \varepsilon_1}{\Gamma \vdash \mathbf{def} \ m(x : \tau_1) : \{\varepsilon_1\} \tau_2 = e : \mathbf{def} \ m(x : \tau_1) : \{\varepsilon_1\} \tau_2} \text{ (DT-DEF)} \\
\\
\frac{\Gamma \vdash x : \{\} \tau}{\Gamma \vdash \mathbf{var} \ f : \tau = x : \mathbf{var} \ f : \tau} \text{ (DT-VAR)} \quad \frac{\Gamma \vdash \varepsilon \text{ wf}}{\Gamma \vdash \mathbf{effect} \ g = \{\varepsilon\} : \mathbf{effect} \ g = \{\varepsilon\}} \text{ (DT-EFFECT)}
\end{array}$$

Figure 3.4: Wyvern static semantics.

is a union between the effect sets the two expressions that are involved in the field assignment produce.

A type substitution of an expression may happen only if the expression is well typed using the original type, the original type is a subtype of the new type, and when the effect set of the original set is a subeffect of the effect of the new type. (Subeffecting is discussed in Section 3.2.5.)

None of the object declarations produce effects, and so object-declaration type-checking rules do not include an effect set preceding the type annotation. For declarations, the judgement reads that, in the variable typing context Γ , the declaration d is a well-typed declaration with the type σ .

When type-checking a method declaration, the effect set annotating the method must be well formed in the overall typing context extended with the method argument. Furthermore, the effect annotating the method must be a supereffect of the effect the method body actually produced.

A field declaration is trivially well typed, and an effect declaration is well typed if the effect set that it is defined with is well formed in the given context.

$$\boxed{\Gamma \vdash \varepsilon <: \varepsilon'}$$

$$\frac{\varepsilon_1 \subseteq \varepsilon_2}{\Gamma \vdash \varepsilon_1 <: \varepsilon_2} \text{ (SUBEFFECT-SUBSET)}$$

$$\frac{\Gamma \vdash n : \{y \Rightarrow \sigma\} \quad \text{effect } g \leq \varepsilon \in \sigma \quad \Gamma \vdash [n/y]\varepsilon \cup \varepsilon_1 <: \varepsilon_2}{\Gamma \vdash \varepsilon_1 \cup \{n.g\} <: \varepsilon_2} \text{ (SUBEFFECT-UPPERBOUND)}$$

$$\frac{\Gamma \vdash n : \{y \Rightarrow \sigma\} \quad \text{effect } g \geq \{\varepsilon\} \in \sigma \quad \Gamma \vdash \varepsilon_1 <: [n/y]\varepsilon \cup \varepsilon_2}{\Gamma \vdash \varepsilon_1 <: \varepsilon_2 \cup \{n.g\}} \text{ (SUBEFFECT-LOWERBOUND)}$$

$$\frac{\Gamma \vdash n : \{y \Rightarrow \sigma\} \quad \text{effect } g = \{\varepsilon\} \in \sigma \quad \Gamma \vdash \varepsilon_1 <: [n/y]\varepsilon \cup \varepsilon_2}{\Gamma \vdash \varepsilon_1 <: \varepsilon_2 \cup \{n.g\}} \text{ (SUBEFFECT-DEF-1)}$$

$$\frac{\Gamma \vdash n : \{y \Rightarrow \sigma\} \quad \text{effect } g = \{\varepsilon\} \in \sigma \quad \Gamma \vdash [n/y]\varepsilon \cup \varepsilon_1 <: \varepsilon_2}{\Gamma \vdash \varepsilon_1 \cup \{n.g\} <: \varepsilon_2} \text{ (SUBEFFECT-DEF-2)}$$

Figure 3.5: Wyvern subeffecting rules.

3.2.5 Subtyping

Subeffecting Rules

As we already saw in the T-SUB, and DT-DEF rules above and as we will see more in the upcoming Section 3.2.5, to compare two sets of effects, we use subeffecting rules, which are presented in Fig. 3.5. If an effect is a subset of another effect, then the former effect is a subeffect of the latter (SUBEFFECT-SUBSET). If an effect set contains an effect variable that is declared with an upper bound, and the union of the rest of the effect set with the upper bound is a subeffect of another effect set, then the former effect set is a subeffect of the latter effect set (SUBEFFECT-UPPERBOUND). If an effect set contains an effect variable that is declared with a lower bound, and the union of the rest of the effect set with the lower bound is a supereffect of another effect set, then the former effect set is a supereffect of the latter (SUBEFFECT-LOWERBOUND). If an effect set contains an effect variable that has a definition, and the union of the rest of the effect set with the definition of the variable is a supereffect of another effect set, then the former effect set is a supereffect of the latter (SUBEFFECT-DEF-1). Finally, if an effect set contains an effect variable that has a definition, and the union of the rest of the effect set with the definition of the variable is a subeffect of another effect set, then the former effect set is a subeffect of the latter (SUBEFFECT-DEF-2).

Lemma 1. $size(\Gamma, \varepsilon)$ (Defined in Fig. 3.6) is finite.

Proof. By rules Safe-1, Safe-2, Safe-3, and Safe-4 in Fig. 3.3, the size of an arbitrary effect $x.g$ is bounded by the total number of effects in the context Γ . \square

Theorem 2. $\Gamma \vdash \varepsilon <: \varepsilon'$ is decidable.

Proof. The proof is by induction on $size(\Gamma, \varepsilon \cup \varepsilon')$.

$$\boxed{size(\Gamma, \varepsilon) = n}$$

$$\frac{}{size(\Gamma, \{\}) = 0} \text{ (SIZE-EMPTY)}$$

$$\frac{\Gamma \vdash x : \{y \Rightarrow \sigma\} \quad \text{effect } g \in \sigma}{size(\Gamma, x.g) = 0} \text{ (SIZE-ABSTRACT)}$$

$$\frac{}{size(\Gamma, \overline{x.g}) = \sum_{x.g \in \overline{x.g}} size(\Gamma, x.g)} \text{ (SIZE-LIST)}$$

$$\frac{\Gamma \vdash x : \{y \Rightarrow \sigma\} \quad \text{effect } g = \{\varepsilon\} \in \sigma}{size(\Gamma, x.g) = 1 + size(\Gamma, [x/y]\varepsilon)} \text{ (SIZE-DEF)}$$

$$\frac{\Gamma \vdash x : \{y \Rightarrow \sigma\} \quad \text{effect } g \leq \{\varepsilon\} \in \sigma}{size(\Gamma, x.g) = 1 + size(\Gamma, [x/y]\varepsilon)} \text{ (SIZE-UPPERBOUND)}$$

$$\frac{\Gamma \vdash x : \{y \Rightarrow \sigma\} \quad \text{effect } g \geq \{\varepsilon\} \in \sigma}{size(\Gamma, x.g) = 1 + size(\Gamma, [x/y]\varepsilon)} \text{ (SIZE-LOWERBOUND)}$$

Figure 3.6: Rules for determining the size of effect definitions.

- BC Since size for both effect is 0, the only applicable rule for subeffecting is Subeffect-Subset. The rule only checks if ε is a subset of ε' , therefore is decidable.
- IS Assume the judgment $\Gamma \vdash \varepsilon <: \varepsilon'$ is derived from Subeffect-Upperbound. In the premise of this rule, we have $\Gamma \vdash [n/y]\varepsilon \cup \varepsilon_1 <: \varepsilon_2$. Since we extract the definition of $n.g$ to find ε , we have $size(\Gamma, [n/y]\varepsilon \cup \varepsilon_1 \cup \varepsilon_2) < size(\Gamma, \{n.g\} \cup \varepsilon_1 \cup \varepsilon_2)$. We can then use induction hypothesis to show the subeffecting judgment in the premise is decidable. The inductive step for rules Subeffect-Lowerbound, Subeffect-Def-1, and Subeffect-Def-2 have the similar structure.

□

Declarative Subtyping Rules

Wyvern subtyping rules are shown in Fig. 3.7. Since, to compare types, we need to compare the effects in them using subeffecting, subtyping relationship is checked in a particular variable typing context. The first four object-subtyping rules and the S-REFL2 rule are standard. In S-DEPTH, since effects may contain a reference to the current object, to check the subtyping relationship between two type declarations, we extend the current typing context with the current object. Method-declaration typing is contravariant in the argument types and covariant in the return type. Furthermore, there must be a covariant-like relationship between the effect sets in the method annotations on the two method declarations: the effect set of the subtype method declaration must be a subeffect of the effect set of the supertype method declaration (S-DEF). An effect definition or an effect declaration with bound is trivially a subtype of an effect declaration (S-EFFECT-1, S-EFFECT-2, S-EFFECT-5). An effect definition is a subtype of an effect declaration with upper bound if the definition is a subeffect of the upper bound (S-EFFECT-3).

$$\boxed{\Gamma \vdash \tau <: \tau'}$$

$$\begin{array}{c} \overline{\Gamma \vdash \tau <: \tau} \text{ (S-REFL1)} \quad \frac{\Gamma \vdash \tau_1 <: \tau_2 \quad \Gamma \vdash \tau_2 <: \tau_3}{\Gamma \vdash \tau_1 <: \tau_3} \text{ (S-TRANS)} \\ \\ \frac{\{x \Rightarrow \sigma_i^{i \in 1..n}\} \text{ is a permutation of } \{x \Rightarrow \sigma_i'^{i \in 1..n}\}}{\Gamma \vdash \{x \Rightarrow \sigma_i^{i \in 1..n}\} <: \{x \Rightarrow \sigma_i'^{i \in 1..n}\}} \text{ (S-PERM)} \\ \\ \overline{\Gamma \vdash \{x \Rightarrow \sigma_i^{i \in 1..n+k}\} <: \{x \Rightarrow \sigma_i^{i \in 1..n}\}} \text{ (S-WIDTH)} \quad \frac{\forall i, \Gamma, x : \{x \Rightarrow \sigma_i^{i \in 1..n}\} \vdash \sigma_i <: \sigma_i'}{\Gamma \vdash \{x \Rightarrow \sigma_i^{i \in 1..n}\} <: \{x \Rightarrow \sigma_i'^{i \in 1..n}\}} \text{ (S-DEPTH)} \end{array}$$

$$\boxed{\Gamma \vdash \sigma <: \sigma'}$$

$$\begin{array}{c} \overline{\Gamma \vdash \sigma <: \sigma} \text{ (S-REFL2)} \quad \frac{\Gamma \vdash \tau'_1 <: \tau_1 \quad \Gamma \vdash \tau_2 <: \tau'_2 \quad \Gamma, x : \tau_1 \vdash \varepsilon_1 <: \varepsilon_2}{\Gamma \vdash \text{def } m(x : \tau_1) : \{\varepsilon_1\} \tau_2 <: \text{def } m(x : \tau'_1) : \{\varepsilon_2\} \tau'_2} \text{ (S-DEF)} \\ \\ \overline{\Gamma \vdash \text{effect } g = \{\varepsilon\} <: \text{effect } g} \text{ (S-EFFECT-1)} \quad \overline{\Gamma \vdash \text{effect } g \leq \varepsilon <: \text{effect } g} \text{ (S-EFFECT-2)} \\ \\ \frac{\Gamma \vdash \varepsilon <: \varepsilon'}{\Gamma \vdash \text{effect } g = \{\varepsilon\} <: \text{effect } g \leq \varepsilon'} \text{ (S-EFFECT-3)} \quad \frac{\Gamma \vdash \varepsilon <: \varepsilon'}{\Gamma \vdash \text{effect } g \leq \varepsilon <: \text{effect } g \leq \varepsilon'} \text{ (S-EFFECT-4)} \\ \\ \overline{\Gamma \vdash \text{effect } g \geq \varepsilon <: \text{effect } g} \text{ (S-EFFECT-5)} \quad \frac{\Gamma \vdash \varepsilon' <: \varepsilon}{\Gamma \vdash \text{effect } g = \{\varepsilon\} <: \text{effect } g \geq \varepsilon'} \text{ (S-EFFECT-6)} \\ \\ \frac{\Gamma \vdash \varepsilon' <: \varepsilon}{\Gamma \vdash \text{effect } g \geq \varepsilon <: \text{effect } g \geq \varepsilon'} \text{ (S-EFFECT-7)} \end{array}$$

Figure 3.7: Wyvern subtyping rules.

Similarly, an effect definition is a subtype of an effect declaration with lower bound if the definition is a supereffect of the lower bound (S-EFFECT-6). An effect declaration with upper bound is a subtype of the effect declaration with another upper bound if the former upper bound is a subeffect of the latter upper bound (S-EFFECT-4). Finally, an effect declaration with lower bound is a subtype of the effect declaration with another lower bound if the former upper bound is a supereffect of the latter upper bound (S-EFFECT-7).

Algorithmic Subtyping Rules

$$\boxed{\Gamma \vdash \tau <: \tau'}$$

$$\frac{\exists \text{ an injection } p : \{1..n\} \mapsto \{1..m\}, \quad \forall i \in 1..n, \Gamma, x : \{x \Rightarrow \sigma_i^{i \in 1..m}\} \vdash \sigma_{p(i)} <: \sigma_i'}{\Gamma \vdash \{x \Rightarrow \sigma_i^{i \in 1..m}\} <: \Gamma \vdash \{x \Rightarrow \sigma_i'^{i \in 1..n}\}} \text{ (S-ALG)}$$

Figure 3.8: Algorithmic Subtyping

The S-Alg rule encodes the S-Refl-1, S-Perm, S-Depth, and S-Width rule using an injective function p . The subtyping rules of declaration types are identical to the declarative subtyping. We prove that S-Trans rules is emissible in theorem 3. Since subtyping rules object types and declaration types are syntax-directed, the subtyping of our effect system is decidable.

Theorem 3. (*Transitivity of algorithmic subtyping*)

If $\Gamma \vdash \tau_1 <: \tau_2$ and $\Gamma \vdash \tau_2 <: \tau_3$, then $\Gamma \vdash \tau_1 <: \tau_3$.

If $\Gamma \vdash \sigma_1 <: \sigma_2$ and $\Gamma \vdash \sigma_2 <: \sigma_3$, then $\Gamma \vdash \sigma_1 <: \sigma_3$.

3.2.6 Dynamic Semantics and Type Soundness

Object-Oriented Core Syntax

$n ::= x \mid l$	<i>names</i>	$\sigma ::= \text{def } m(x : \tau) : \{\varepsilon\} \tau$	<i>declaration types</i>
$e ::= n$	<i>expressions</i>	$\text{var } f : \tau$	
$\quad \mid \text{new}(x \Rightarrow \bar{d})$		$\text{effect } g$	
$\quad \mid e.m(e)$		$\text{effect } g \geq \{\varepsilon\}$	
$\quad \mid e.f$		$\text{effect } g \leq \{\varepsilon\}$	
$\quad \mid e.f = e$		$\text{effect } g = \{\varepsilon\}$	
$\varepsilon ::= \bar{n}.\bar{g}$	<i>effects</i>	$\Gamma ::= \emptyset \mid \Gamma, x : \tau$	<i>var. typing context</i>
$d ::= \text{def } m(x : \tau) : \{\varepsilon\} \tau = e$	<i>declarations</i>	$\mu ::= \emptyset \mid \mu, l \mapsto \{x \Rightarrow \bar{d}\}$	<i>store</i>
$\quad \mid \text{var } f : \tau = n$		$\Sigma ::= \emptyset \mid \Sigma, l : \tau$	<i>store typing context</i>
$\quad \mid \text{effect } g = \{\varepsilon\}$		$E ::= []$	<i>evaluation context</i>
$\tau ::= \{x \Rightarrow \bar{\sigma}\}$	<i>object type</i>	$E.m(e)$	
		$l.m(E)$	
		$E.f$	
		$E.f = e$	
		$l.f = E$	

Figure 3.9: Wyvern’s object-oriented core syntax with dynamic forms.

Fig. 3.9 shows the version of the syntax of Wyvern’s object-oriented core that includes dynamic semantics. Specifically, expressions include locations l , which variables in effects resolve to at run time. We also use a store μ and its typing context Σ . Finally, to make the dynamics more compact we use an evaluation context E .

Changes in Static Semantics

Type checking a location (T-LOC) and a field declaration (DT-VAR) is straightforward, and we also need to ensure that the store is well-formed and contains objects that respect their types.

Dynamic Semantics

The dynamic semantics that we use for Wyvern’s effect system is shown in Fig. 3.11 and is similar to the one described in prior work [19]. In comparison to the prior work, this version of Wyvern’s dynamic semantics has fewer rules, and the E-METHOD rule is simplified.

$$\begin{array}{c}
\boxed{\Gamma \mid \Sigma \vdash e : \{\varepsilon\} \tau} \\
\ldots \quad \frac{l : \tau \in \Sigma}{\Gamma \mid \Sigma \vdash l : \{\} \tau} \text{ (T-LOC)} \\
\\
\boxed{\Gamma \mid \Sigma \vdash d : \sigma} \\
\ldots \quad \frac{\Gamma \mid \Sigma \vdash n : \{\} \tau}{\Gamma \mid \Sigma \vdash \text{var } f : \tau = n : \text{var } f : \tau} \text{ (DT-VAR)} \\
\\
\boxed{\mu : \Sigma} \\
\frac{\forall l \mapsto \{x \Rightarrow \bar{d}\} \in \mu, \forall i, d_i \in \bar{d}, \sigma_i \in \bar{\sigma}, x : \{x \Rightarrow \bar{\sigma}\} \mid \Sigma \vdash d_i : \sigma_i}{\mu : \Sigma} \text{ (T-STORE)}
\end{array}$$

Figure 3.10: Wyvern static semantics affected by dynamic semantics.

The judgement reads the same as before: given the store μ , the expression e evaluates to the expression e' and the store becomes μ' . The E-CONGRUENCE rule still handles all non-terminal forms. To create a new object (E-NEW), we select a fresh location in the store and assign the object's definition to it. Provided that there is an appropriate method definition in the object on which a method is called, the method call is reduced to the method's body (E-METHOD). In the method's body, the locations representing the method argument and the object on which the method is called are substituted for corresponding variables. An object field is reduced to the value held in it (E-FIELD), and when an object field's value changes (E-ASSIGN), appropriate substitutions are made in the object's declaration set and the store.

Type Soundness

We prove the soundness of the effect system presented above using the standard combination of progress and preservation theorems. Proof to these theorems can be found in Appendix ??.

Theorem 4 (Preservation). *If $\Gamma \mid \Sigma \vdash e : \{\varepsilon\} \tau$, $\mu : \Sigma$, and $\langle e \mid \mu \rangle \longrightarrow \langle e' \mid \mu' \rangle$, then $\exists \Sigma' \supseteq \Sigma$, $\mu' : \Sigma'$, $\exists \varepsilon'$, such that $\Gamma \vdash \varepsilon' <: \varepsilon$, and $\Gamma \mid \Sigma' \vdash e' : \{\varepsilon'\} \tau$.*

Theorem 5 (Progress). *If $\emptyset \mid \Sigma \vdash e : \{\varepsilon\} \tau$ (i.e., e is a closed, well-typed expression), then either*

1. *e is a value (i.e., a location) or*
2. *$\forall \mu$ such that $\mu : \Sigma$, $\exists e', \mu'$ such that $\langle e \mid \mu \rangle \longrightarrow \langle e' \mid \mu' \rangle$.*

$$\boxed{\langle e \mid \mu \rangle \longrightarrow \langle e' \mid \mu' \rangle}$$

$$\frac{\langle e \mid \mu \rangle \longrightarrow \langle e' \mid \mu' \rangle}{\langle E[e] \mid \mu \rangle \longrightarrow \langle E[e'] \mid \mu' \rangle} \text{ (E-CONGRUENCE)} \qquad \frac{l \notin \text{dom}(\mu)}{\langle \text{new}(x \Rightarrow \bar{d}) \mid \mu \rangle \longrightarrow \langle l \mid \mu, l \mapsto \{x \Rightarrow \bar{d}\} \rangle} \text{ (E-NEW)}$$

$$\frac{l_1 \mapsto \{x \Rightarrow \bar{d}\} \in \mu \quad \text{def } m(y : \tau_1) : \{\varepsilon\} \tau_2 = e \in \bar{d}}{\langle l_1.m(l_2) \mid \mu \rangle \longrightarrow \langle [l_2/y][l_1/x]e \mid \mu \rangle} \text{ (E-METHOD)}$$

$$\frac{l \mapsto \{x \Rightarrow \bar{d}\} \in \mu \quad \text{var } f : \tau = l_1 \in \bar{d}}{\langle l.f \mid \mu \rangle \longrightarrow \langle l_1 \mid \mu \rangle} \text{ (E-FIELD)}$$

$$\frac{l_1 \mapsto \{x \Rightarrow \bar{d}\} \in \mu \quad \text{var } f : \tau = l \in \bar{d} \quad \bar{d}' = [\text{var } f : \tau = l_2 / \text{var } f : \tau = l] \bar{d} \quad \mu' = [l_1 \mapsto \{x \Rightarrow \bar{d}'\} / l_1 \mapsto \{x \Rightarrow \bar{d}\}] \mu}{\langle l_1.f = l_2 \mid \mu \rangle \longrightarrow \langle l_2 \mid \mu' \rangle} \text{ (E-ASSIGN)}$$

Figure 3.11: Wyvern dynamic semantics.

Chapter 4

Abstract Algebraic Effects Via Embedding

4.1 Background and Motivation

Algebraic effects (introduced by Plotkin and Power [24]) and handlers (introduced by Plotkin and Pretnar [26]) are an approach to computational effects based on a premise that impure behavior arises from a set of operations, and are recently gaining popularity due to its ability to model various form of computational effects such as exceptions, mutable states, async-await, etc.

Modularity is a key concept that separates abstract algebraic effects from the traditional way of using monad to model effects in purely functional programming [28]. [Insert Reasons Here]

However, similar to the restrictive strand of work on effects, few works on algebraic effects has investigated the algebraic effect system on a larger scale, where abstraction between program components is important. Biernacki et al. [2] first studied abstraction of algebraic effects and handlers, presented a calculus equipped with existential algebraic effects and typed runtime coercions, and provided an implementation as a proof of concept.

Abstract algebraic effects are first introduced in [2]. Similar to abstract types, abstract effect signature allows program components to define abstract effect signatures that is opaque to other components in the system. The difference between concrete and abstract effect signatures lies in the ability for program components to handle them. If an effect signature is concrete to a program component, then the operations are accessible the component, and a handler can handles the effect by handling the operations in the signature. On the other hand, if an effect signature is abstract to one program component, then the component should not observe the operations defined in the effect signature, and is therefore unable to handle the effect. As abstraction is an important issue for module systems because it provides a separation of implementation details of functions from the interface, abstraction of algebraic effects provides a similar benefit for modularity because it helps separate the component operations from the effect signature, ensuring that the client can only uses the handler provided by the library to handle the effect.

The following code is a motivating example similar to the example in [2] that illustrates the challenges of implementing abstract algebraic effects. `Nondet` is a globally defined effect signature. Then we define a module `m` with type `M` with an abstract effect `E`, a method `mflip`, and a handler method `handle`. The effect `E` is defined by `Nondet`, but is opaque to the outside world of the module, because `E` is defined as an abstract effect in the type `M`. The method `mflip` simply

calls the `flip` operation, and the `handle` method handles the `flip` operation by returning `true`.

```

1 effect Nondet {
2     flip(): Bool
3 }
4
5 type M
6     effect E
7     def mflip() : {this.E} Bool
8     def handler(Unit -> {this.E} Bool) : {} Bool
9
10 module m: M
11     effect E = {Nondet}
12     def mflip() : {this.E} Unit
13         flip()
14     def handle(c: Unit -> {this.E} Bool) : {} Bool =
15         handle c() with
16         | flip() -> return true
17
18 m.handle(
19     () => handle m.mflip() with
20         | flip() -> return false
21 )

```

The last segment of the above example shows an client code of module `m` that calls the method `m.handle` and pass in an expression that encapsulate the call to `m.mflip` by another handler that handles the `flip` operation. Since the effect of the method `m.mflip` is abstract, the inner handler should not handle the operation inside `m.mflip`. Instead, the operation should be handled by the outer handler method `m.handle`. However,

As we can see, the abstraction of effect signatures from type abstractions, since the erasure of type information would make the abstraction unsound. So we need a language that keep track of the information on effect abstraction during the evaluation of the program. In this work, we incorporate the method of syntactic type abstraction introduced by Grossman et al. [10], who use the notion of principals to track the flow of values with abstract types during the evaluation of a program.

4.2 Core Calculus

4.2.1 Syntax

This section describes a variant of the simply typed lambda calculus that maintains a syntactic distinction between agents during evaluation. Figure 5.3.1 gives the syntax of our calculus. As our previous discussion, it is crucial to keep track of the effect abstraction information during the evaluation of the program. It is therefore natural to divide the code into agents, and allow each agent to export abstract effect signatures. We assume that there are n agents, and use variables i, j, k to range over the set of agents.

Every term in this language is assigned to an agent. And terms are split into inert expressions and potentially effectful computations, following an approach called *fine-grain call-by-value*,

(agents)	$i, j ::= \{1 \dots n\}$
(lists)	$l ::= i \mid il$
(value types)	$\tau ::= \text{unit} \mid \tau \rightarrow \sigma$
(computation types)	$\sigma ::= \{\varepsilon\}\tau$
(effect types)	$\varepsilon ::= \cdot \mid f, \varepsilon \mid op, \varepsilon$
(i values)	$v_i ::= ()_i \mid \lambda x_i : \tau. c_i$
(i expression)	$e_i ::= x_i \mid v_i \mid [e_j]_j^\tau$
(i computation)	$c_i ::= \text{return } e_i \mid op(e_i, y.c_i) \mid \text{do } x \leftarrow c_i \text{ in } c'_i \mid e_i e'_i \mid \text{with } h_i \text{ handle } c_i$ $\mid [c_j]_j^\sigma \mid [op]_l^\varepsilon(e_i, y_i.c_i)$
(i handler)	$h_i ::= \text{handler } \{\text{return } x_i \mapsto c_i^r, op^1(x_i^1, k^1) \mapsto c_i^1 \dots op^n(x_i^n, k^n) \mapsto c_i^n\}$

Figure 4.1: Syntax for multi-agent calculus

introduced by Levy et al. [15]. We use the notation *i-expression* and *i-computation* to denote expressions and computations in the agent *i*. We use subscripts to indicate a term is assigned to an agent, however, we will omit the subscript if the agent the term belongs to is not important or obvious in the context.

An *i-value* is an *i-expression* that cannot further reduce. There are two forms of *i-value*: the unit $()$, and the lambda abstraction $\lambda x_i : \tau. c_i$. *i-expressions* include variable x_i , value v_i , and embedded expressions $[e_j]_j^\tau$. *i-computations* are the terms that can potentially cause effects, and consists of return statement $\text{return } e_i$, operation call $op(e_i; y_i.c_i)$, sequencing $\text{do } x_i \leftarrow c_i \text{ in } c'_i$, application $e_i e'_i$, handling $\text{with } h_i \text{ handle } c_i$, embedded computation $[c_j]_j^\sigma$, and embedded operation call $[op]_l^\varepsilon(e_i; y_i.c_i)$. There are few things worth mentioning:

Sequencing: in $\text{do } x \leftarrow c \text{ in } c'$, we first evaluate c , bind the return value of c' to x and then evaluate c_2

Operation Calls: The call $op(e; y.c)$ passes the parameter e to the operation op , binds the return value of the operation call to y , and continue by evaluating the computation c . Note that the encompassing handler could potentially change the behavior of the operation.

Embeddings: the term $[e_j]_j^\tau$ is *i-expression* that contains an embedded *j-expression*, with a type τ exported by j . Similarly, $[c_j]_j^\sigma$ is an embedded *j-expression* with exported type σ

Embedded Operations: The embedded operation $[op]_l^\varepsilon(e; y.c)$ is a operation call that is annotated with effect ε . l is a list of agents that have contributed to the formation of the annotation. We will describe this in more detail later.

Similar to terms, types are also divided into expression types and computation types. There are two forms of expression types τ : the unit type 1 , and the arrow type $\tau \rightarrow \sigma$. As for the computation type σ , there is only one form: $\{\varepsilon\}\tau$: where ε is a set of effect that the computation might induce, and τ is the type of the return value of the computation.

The effect type ε represents an unordered set of effects that can be empty \cdot . A effect type can be extended by either an effect label f , or an operation op .

The i -handler h_i must contain a return clause that handles the case when the handled computation directly returns a value. It may also contain clauses that handle operations. For example, the clause $op(x, k) \mapsto c$ handles the operation op . More details can be found in the dynamic semantics section.

4.2.2 Agent Specific Type Information

We use agents to model a module system where each module can have private information about effect abstraction. Each agent in our language has limited knowledge of effect abstraction. For example, an agent i might know that `effect Nondet = flip() : Bool`, and an agent j does not have this information. As a result, agent i would be able to handle a computation with effect `Nondet`, while the agent j would not be able to do that. Furthermore, we need to ensure the consistency of the information on effect abstraction, that is, agent j should not think that the effect `Nondet = read() : String`, which would contradict with the knowledge of agent i .

The model of effect abstraction information is similar to the model of type information in [10]. To capture effect abstraction information, each agent i has a partial function δ_i that maps an effect label to an effect type. There are two requirements for these maps: (1) For each effect label f , if there are two agents that know the implementation of the effect f , then their knowledge about the implementation must be the same. (2) For each effect label f , there is a unique and most concrete interpretation of f . We would not allow the effect label f itself to appear in the implementation of f . Examples like $\delta_i(f) = \{f\}$ would be rejected.

Definition 4.2.1. A set $\{\delta_1, \dots, \delta_n\}$ of maps from effect labels to effects is compatible if

1. For all $i, j \in 1 \dots n$ if $f \in \text{Dom}(\delta_i) \cap \text{Dom}(\delta_j)$, then $\delta_i(f) = \delta_j(f)$.
2. Effect labels can be totally ordered such that for every agent i and effect label f , all effect labels in $\delta_i(f)$ precede f .

Then we define the total function Δ_i that refines an effect type:

Definition 4.2.2.

$$\begin{aligned} \Delta_i(\cdot) &= \cdot \\ \Delta_i(op, \varepsilon) &= op, \Delta_i(\varepsilon) \\ \Delta_i(f, \varepsilon) &= \begin{cases} f, \Delta_i(\varepsilon) & \text{if } f \notin \text{Dom}(\delta_i) \\ \varepsilon', \Delta_i(\varepsilon) & \text{if } \delta_i(f) = \varepsilon' \end{cases} \end{aligned}$$

The definition of compatibility ensures that there is a fixpoint for repeatedly refining an effect label f using the function Δ_i . We call such fixpoint $\overline{\Delta_i}(f)$.

Definition 4.2.3. $\overline{\Delta_i}(f) = \varepsilon$ if there is some $n \geq 0$

$$\underbrace{\Delta_i(\dots(\Delta_i(f)\dots))}_{n \text{ applications}} = \underbrace{\Delta_i(\dots(\Delta_i(f)\dots))}_{n+1 \text{ applications}} = \varepsilon$$

We assume that the type information for operations are public to all agents. The type for an operation op is contained in a separate map Σ , which maps an operation op to an arrow type $\tau_A \rightarrow \tau_B$. Note that this is different from the function type in our calculus, which has the form $\tau \rightarrow \sigma$.

$$\boxed{e \longrightarrow e'}$$

$$\frac{e_j \mapsto e'_j}{[e_j]_j^\tau \longrightarrow [e'_j]_j^\tau} \text{ (E-CONGRUENCE)} \quad \frac{}{[(\cdot)]_j^1 \longrightarrow (\cdot)_i} \text{ (E-UNIT)}$$

$$\frac{}{[\lambda x_j : \tau'. c_j]_j^{\tau \rightarrow \sigma} \longrightarrow \lambda x_i : \tau. [\{[x_i]_i^{\tau'} / x_j\} c_j]_j^\sigma} \text{ (E-LAMBDA)}$$

Figure 4.2: Operational Semantics for Expressions

4.3 Operational Semantics

The reduction rules for terms are dependent on the agent of the terms. Figure 4.2 shows that operational semantics for expressions of agent i . (E-Congruence) shows that a j -expression embedded agent- i should be evaluated using the reductions rules for agent j first. The (E-Unit) and (E-Lambda) rules show that we can lift an embedded j -value to agent i , so the value becomes an i -value. The (E-Unit) rule simply lifts the unit value out of the embedding. The (E-Lambda) rule is more interesting: The value embedded is a lambda expression of agent j . We lift the argument out of the embedding. However, the type annotating the argument is changed from τ' to the exported argument type τ , because the reduced expression should have the exported type $\tau \rightarrow \sigma$. The body of the reduced expression is an embedded j -computation, so the variable x_i should be encapsulated by an embedding, because any i -term should be embedded in a j -term. We annotate x_i with type τ' because the original lambda function expects a value of type τ' .

Figure 4.3 shows the reduction rules for i -computations. (E-Ret) is the congruence rule that evaluates the expression in return statement. (E-Op) evaluates the input argument for the operation call. Note that there are no reduction rules for operation calls because the semantics for operations are defined by the handler encapsulating it.

(E-EmbedOp1) evaluates the argument expression to an embedded operation. (E-EmbedOp2) refines the effect annotation of an operation. (E-EmbedOp3) lifts the operation out of an embedding when the annotation contains the operation, because the agent i has enough information about effect abstraction to handle the operation. Note that we require that the effect annotation cannot be further refined, in order to ensure determinism of evaluation. (E-EmbedOp4) removes an operation that is not op out of the effect annotation, as this is helpful in our proof of type soundness.

(E-App1), (E-App2) and (E-App3) are standard call-by-value semantics for applications. (E-Seq1) evaluates the first computation in a sequence of computations. (E-Seq2) binds the return value of the first computation to a variable in the second computation. (E-Seq3) witnesses an operation call as the first computation in a sequence. Since there is no way to further evaluate an operation right away, we propagate the operation call outwards and defer further evaluation to the continuation of the call. (E-Seq4) is similar to (E-Seq3), and requires that the effect annotation on the embedding to be most precise.

(E-Handle1) simply evaluates the computation encapsulated by the handler. In (E-Handle2), the computation returns a value, so we substitute the value into the computation of the clause that handles the return statement in the handler. (E-Handle3) shows that case when the handler h_i has

$$\boxed{c \longrightarrow c'}$$

$$\begin{array}{c}
\frac{e_i \mapsto e'_i}{\text{return } e_i \longrightarrow \text{return } e'_i} \text{ (E-RET)} \quad \frac{e_i \mapsto e'_i}{\text{op}(e_i, y_i, c_i) \longrightarrow \text{op}(e'_i, y_i, c_i)} \text{ (E-OP)} \\
\\
\frac{e_i \longrightarrow e'_i}{[\text{op}]_l^\varepsilon(e_i; y_i, c_i) \longrightarrow [\text{op}]_l^\varepsilon(e'_i; y_i, c_i)} \text{ (E-EMBEDOP1)} \quad \frac{\overline{\Delta_i(\varepsilon)} = \varepsilon'}{[\text{op}]_l^\varepsilon(v_i; y_i, c_i) \longrightarrow [\text{op}]_l^{\varepsilon'}(v_i; y_i, c_i)} \text{ (E-EMBEDOP2)} \\
\\
\frac{\Delta_i(\varepsilon) = \varepsilon \quad \text{op} \in \varepsilon}{[\text{op}]_l^\varepsilon(v_i; y_i, c_i) \longrightarrow \text{op}(v_i; y_i, c_i)} \text{ (E-EMBEDOP3)} \quad \frac{\Delta_i(\varepsilon) = \varepsilon \quad \text{op} \notin \varepsilon \quad \text{op}' \in \varepsilon}{[\text{op}]_l^\varepsilon(v_i; y_i, c_i) \longrightarrow [\text{op}]^{\varepsilon \setminus \text{op}'}(v_i; y_i, c_i)} \text{ (E-EMBEDOP4)} \\
\\
\frac{e_i \longrightarrow e''_i}{e_i \ e'_i \longrightarrow e''_i \ e'_i} \text{ (E-APP1)} \quad \frac{e_i \longrightarrow e'_i}{v_i \ e_i \longrightarrow v_i \ e'_i} \text{ (E-APP2)} \quad \frac{}{(\lambda x_i : \tau. c_i) v_i \longrightarrow \{v_i/x_i\}c_i} \text{ (E-APP3)} \\
\\
\frac{c_i \longrightarrow c'_i}{\text{do } x \leftarrow c_i \text{ in } c'_i \longrightarrow \text{do } x \leftarrow c'_i \text{ in } c'_i} \text{ (E-SEQ1)} \quad \frac{}{\text{do } x \leftarrow \text{return } v_i \text{ in } c'_i \longrightarrow \{v_i/x\}c'_i} \text{ (E-SEQ2)} \\
\\
\frac{}{\text{do } x \leftarrow \text{op}_i(v_i; y_i, c_i) \text{ in } c'_i \longrightarrow \text{op}_i(v_i; y_i, \text{do } x \leftarrow c_i \text{ in } c'_i)} \text{ (E-SEQ3)} \\
\\
\frac{\Delta_i(\varepsilon) = \varepsilon \quad \text{op} \notin \varepsilon}{\text{do } x \leftarrow [\text{op}_j]_l^\varepsilon(v_i; y_i, c_i) \text{ in } c'_i \longrightarrow [\text{op}_j]_l^\varepsilon(v_i; y_i, \text{do } x \leftarrow c_i \text{ in } c'_i)} \text{ (E-SEQ4)} \\
\\
\frac{c_i \longrightarrow c'_i}{\text{with } h_i \text{ handle } c_i \longrightarrow \text{with } h_i \text{ handle } c'_i} \text{ (E-HANDLE1)} \quad \frac{\text{return } x_i \mapsto c'_i \in h_i}{\text{with } h_i \text{ handle return } v_i \longrightarrow \{v_i/x_i\}c'_i} \text{ (E-HANDLE2)} \\
\\
\frac{\text{op}(x_i; k) \mapsto c'_i \in h_i \quad \Sigma(\text{op}) = \tau_A \rightarrow \tau_B}{\text{with } h_i \text{ handle } \text{op}(v_i; y_i, c_i) \longrightarrow \{v_i/x_i\}\{(\lambda y : \tau_B. \text{with } h_i \text{ handle } c_i)/k\}c'_i} \text{ (E-HANDLE3)} \\
\\
\frac{\text{op}(x_i; k) \mapsto c'_i \notin h_i}{\text{with } h_i \text{ handle } \text{op}(v_i; y_i, c_i) \longrightarrow \text{op}(v_i; y_i, \text{with } h_i \text{ handle } c_i)} \text{ (E-HANDLE4)} \\
\\
\frac{\Delta_i(\varepsilon) = \varepsilon \quad \text{op} \notin \varepsilon}{\text{with } h_i \text{ handle } [\text{op}]_l^\varepsilon(v_i, y_i, c_i) \longrightarrow [\text{op}]_l^\varepsilon(v_i; y_i, \text{with } h_i \text{ handle } c_i)} \text{ (E-HANDLE5)} \\
\\
\frac{c_j \longrightarrow c'_j}{[c_j]_l^\sigma \longrightarrow [c'_j]_l^\sigma} \text{ (E-EMBED1)} \quad \frac{}{[\text{return } v_j]_l^{\{\varepsilon\}\tau} \longrightarrow \text{return } [v_j]_l^\tau} \text{ (E-EMBED2)} \\
\\
\frac{\Sigma(\text{op}) = \tau_A \rightarrow \tau_B}{[\text{op}_j(v_j; y_j, c_j)]_j^{\{\varepsilon\}\tau} \longrightarrow [\text{op}_j]_l^\varepsilon([v_j]_j^{\tau_A}; y_i, \{[y_i]_i^{\tau_B}/y_j\}[c_j]_j^{\{\varepsilon\}\tau})} \text{ (E-EMBED3)} \\
\\
\frac{\Sigma(\text{op}) = \tau_A \rightarrow \tau_B \quad \Delta_j(\varepsilon') = \varepsilon' \quad \text{op} \notin \varepsilon'}{[[\text{op}_k]_l^{\varepsilon'}(v_j; y_j, c_j)]_j^{\{\varepsilon\}\tau} \longrightarrow [\text{op}_k]_l^\varepsilon([v_j]_j^{\tau_A}; y_i, \{[y_i]_i^{\tau_B}/y_j\}[c_j]_j^{\{\varepsilon\}\tau})} \text{ (E-EMBED4)}
\end{array}$$

Figure 4.3: Operational Semantics for Computations

a matching clause for the operation op . We substitute the argument v_i for x_i , and substitute the continuation of the operation for k . The continuation function receives an argument of type τ_B , which is the result type of the operation op , and computes the continuation of the operation c_i , but encapsulating it with the handler h_i . The (E-Handle4) shows the case when the operation is not handled by the handler, so we propagate the operation outwards to wait for another handler to handle it. The (E-Handler5) ensures that abstracted effects are not handled: If the current agent cannot refine the effect annotation, then the operation is abstract and cannot be handled, and is therefore propagated outward.

(E-Embed1) is the congruence rule for embedded computations. (E-Embed2) lifts the return statement out of the embedding. We can safely remove the effect annotation ε because the returned value v_j cannot cause any effect.

(E-Embed3) lifts an operation call out of the embedding. We annotate the operation with the effect annotation of the whole computation. Since the argument value for op is a j-value, we need to embed it as a i-value, and annotate it with type τ_A . The continuation c_j is still embedded, and we substitute the embedded variable y_i for y_j , because y_i should be an embedded i-value in a j-value.

The (E-Embed4) rule is very similar to (E-Embed3) with one difference being that op is already embedded. In this case, we override the annotation on op with the annotation for the whole computation, and update the agent list in the subscript of the operation. We add j to the agent list because the agent j have contributed to the effect annotation of the operation.

4.4 Static Semantics

4.4.1 Typing Rules

Figure 4.5 shows the static semantics of the core-calculus. Static semantics contains typing rules for both expressions and computations. Note that the type rules depend on the agent each expression or computation belongs to.

The rule (T-Unit) assigns the unit type to a unit value. (T-Var) looks up a type of a variable from the context. (T-Lam) is the standard rule for typing a lambda function. Note that the body of a lambda is computation, so we need use the typing judgement for computation in the premise of this rule. (T-EmbedExp) assigns type to expression embeddings: The embedding has type τ_{au} if the embedded expression e_j is assigned to the type τ' , and τ_{au} is related to τ_{au}' by the list li . We will elaborate on type relations later.

(T-Ret) assign a type to a return statement: As expected, the expression part of the computation type matches the type of the returned expression. However we can annotate the return statement with an arbitrary effect set, because our type system does not describe the precise effect in computations, but gives the upper bound of effect in computations.

(T-Op) shows the typing rule for operation calls. Again, since we allow effect annotations to be an unupper bound on effect, we can require the operation op to be in the effect set ε .

(T-Handle) shows the typing rule for the effect handling statement with h_i handle c_i . c_i is a computation with effect type ε and return type τ_A . h_i is a handler that contains clause that handler operations op^1, \dots, op^n . For the return clause, given the type of variable x is τ_A , the type

$$\boxed{\Gamma \vdash e_i : \tau}$$

$$\begin{array}{c} \overline{\Gamma \vdash ()_i : 1} \text{ (T-UNIT)} \quad \overline{\Gamma \vdash x_i : \Gamma(x_i)} \text{ (T-VAR)} \quad \frac{\Gamma, x_i : \tau \vdash c_i : \sigma}{\Gamma \vdash (\lambda x_i : \tau. c_i) : \tau \rightarrow \sigma} \text{ (T-LAM)} \\ \\ \frac{\Gamma \vdash e_j : \tau' \quad \Gamma \vdash \tau' \leq_{ji} \tau}{\Gamma \vdash [e_j]_j^\tau : \tau} \text{ (T-EMBEDEXP)} \end{array}$$

$$\boxed{\Gamma \vdash c_i : \sigma}$$

$$\begin{array}{c} \frac{\Gamma \vdash e_i : \tau \quad \Delta_i(\varepsilon) = \varepsilon}{\Gamma \vdash \text{return } e_i : \{\varepsilon\}\tau} \text{ (T-RET)} \quad \frac{\Sigma(op) = \tau_A \rightarrow \tau_B \quad \Gamma \vdash e_i : \tau_A \quad \Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon\}\tau \quad op \in \Delta_i(\varepsilon)}{\Gamma \vdash op(e_i; y_i. c_i) : \{\varepsilon\}\tau} \text{ (T-OP)} \\ \\ \frac{\Gamma \vdash c_i : \{\varepsilon\}\tau \quad \Gamma, x_i : \tau \vdash c'_i : \{\varepsilon\}\tau'}{\Gamma \vdash \text{do } x_i \leftarrow c_i \text{ in } c'_i : \{\varepsilon\}\tau'} \text{ (T-SEQ)} \quad \frac{\Gamma \vdash e_1 : \tau \rightarrow \sigma \quad \Gamma \vdash e_2 : \tau}{\Gamma \vdash e_1 e_2 : \sigma} \text{ (T-APP)} \\ \\ \frac{\begin{array}{l} h_i = \{ \text{return } x \mapsto c^r, op^1(x; k) \mapsto c^1, \dots, op^n(x; k) \mapsto c^n \} \\ \Gamma, x : \tau_A \vdash c^r : \{\varepsilon'\}\tau_B \quad \{ \Sigma(op^i) = \tau_i \rightarrow \tau'_i \mid \Gamma, x : \tau_i, k : \tau'_i \rightarrow \{\varepsilon'\}\tau_B \vdash c^i : \{\varepsilon'\}\tau_B \}_{1 \leq i \leq n} \\ \Gamma \vdash c_i : \{\varepsilon\}\tau_A \quad \varepsilon \setminus \{op^i\}_{1 \leq i \leq n} \subseteq \varepsilon' \end{array}}{\Gamma \vdash \text{with } h_i \text{ handle } c_i : \{\varepsilon'\}\tau_B} \text{ (T-HANDLE)} \\ \\ \frac{\Gamma \vdash c_j : \sigma' \quad \Gamma \vdash \sigma' \leq_{li} \sigma}{\Gamma \vdash [c_j]_l^\sigma : \sigma} \text{ (T-EMBED)} \\ \\ \frac{\Sigma(op) = \tau_A \rightarrow \tau_B \quad \Gamma \vdash e_i : \tau_A \quad \Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon\}\tau \quad \overline{\Delta_i(\varepsilon)} \subseteq \overline{\Delta_i(\varepsilon')} \quad \Gamma \vdash op \leq_{li} \varepsilon}{\Gamma \vdash [op]_l^\varepsilon(e_i; y_i. c_i) : \{\overline{\Delta_i(\varepsilon')}\}\tau} \text{ (T-EMBEDOP)} \end{array}$$

Figure 4.4: Static Semantics

of c^r must be $\{\varepsilon'\}\tau_B$. For the clause handling the operation op^i , which has the operation type $\tau_i \rightarrow \tau'_i$, if variable x has type τ_i , and continuation has the type $\tau'_i \rightarrow \{\varepsilon'\}\tau_B$, then the handling computation c^i must have the type $\{\varepsilon'\}\tau_B$. The effect type after handling, ε' should contain all of the effects that are not handled by the handler.

(T-Embed) is very similar to (T-EmbedExp), where we use the type of the embedded computation and the type relation judgment to assign a type to the embedding.

(T-EmbedOp) first computes the type of c_i given that y_i has the correct type. It is required that the effect annotation on the embedding is a part of the effect of c_i . And op should be related to the effect annotation using the type relations.

4.4.2 Type Relations

$$\boxed{\tau \leq_l \tau'}$$

$$\frac{}{1 \leq_l 1} \text{ (R-UNIT)} \quad \frac{\tau \leq_l \tau' \quad \sigma \leq_l \sigma'}{\tau \rightarrow \sigma \leq_l \tau' \rightarrow \sigma'} \text{ (R-ARROW)}$$

$$\boxed{\sigma \leq_l \sigma'}$$

$$\frac{\varepsilon \leq_l \varepsilon' \quad \tau \leq_l \tau'}{\{\varepsilon\}\tau \leq_l \{\varepsilon'\}\tau'} \text{ (R-SIGMA)}$$

$$\boxed{\varepsilon \leq_l \varepsilon'}$$

$$\frac{\overline{\Delta}_i(\varepsilon) \subseteq \overline{\Delta}_i(\varepsilon')}{\varepsilon \leq_i \varepsilon'} \text{ (R-EFF1)} \quad \frac{\varepsilon \leq_l \varepsilon'' \quad \varepsilon'' \leq_{l'} \varepsilon'}{\varepsilon \leq_{ll'} \varepsilon'} \text{ (R-EFF2)}$$

Figure 4.5: Type Relations

Type relations ensures the soundness of abstraction. The goal of type relations is to prohibit embeddings from exporting incorrect effect abstractions. For example, if a i -computation uses an effect operation $flip : 1 \rightarrow bool$, it should annotate the computation with effect $Nondet$ when exporting the computation, but should not annotate it with the empty effect.

The judgements for expression types are of the form $\tau \leq_l \tau$, where l is a list of agents that provide the type abstraction information used by the relation. (R-Unit) shows that unit types relates to itself. (R-Arrow) relates two arrow types given that the input types and the output types are related, To relate two computation types, we just need to ensure the effect types and return types are related.

The relation for effect types does the actual work. By (R-EFF1), two effect types are related under a single agent i if the first effect is a subset of the second effect after refined by the type information provided by i . (R-EFF2) shows that by using type information from a list of agents, we can combine the chain of relation between effects.

4.5 Safety Properties

In this section we state and prove the standard type-safety theorems for the core calculus.

Lemma 6. (*Substitution*)

If $\Gamma, x_j : \tau' \vdash c_i : \sigma$ and $\Gamma \vdash e_j : \tau'$, then $\Gamma \vdash \{e_j/x_j\}c_i : \sigma$, and

If $\Gamma, x_j : \tau' \vdash e_i : \tau$ and $\Gamma \vdash e_j : \tau'$, then $\Gamma \vdash \{e_j/x_j\}e_i : \tau$

Proof. By rule induction on $\Gamma \vdash e : \tau$ and $\Gamma \vdash c : \sigma$

(T-Unit) Trivial

(T-Var) Trivial

(T-Lam)

$$\frac{\Gamma, x_j : \tau', x_i : \tau \vdash c_i : \sigma}{\Gamma, x_j : \tau' \vdash (\lambda x_i : \tau. c_i) : \tau \rightarrow \sigma} \text{ (T-LAM)}$$

By IH, we have $\Gamma, x_i : \tau \vdash \{e_j/x_j\}c_i : \sigma$ Then by (T-Lam) we have $\Gamma \vdash (\lambda x_i : \tau. \{e_j/x_j\}c_i) : \sigma$.Which is equivalent to $\Gamma \vdash \{e_j/x_j\}(\lambda x_i : \tau. c_i) : \sigma$.

(T-EmbedExp) By inversion and IH

(T-Ret) Follows by induction hypothesis

(T-Op)

$$\frac{\Sigma(op) = \tau_A \rightarrow \tau_B \quad \Gamma \vdash e_i : \tau_A \quad \Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon\}\tau \quad op \in \Delta_i(\varepsilon)}{\Gamma \vdash op(e_i; y_i.c_i) : \{\varepsilon\}\tau} \text{ (T-OP)}$$

By inversion we have $\Gamma, x_j : \tau' \vdash e_i : \tau_A$ and $\Gamma, x_j : \tau', y_i : \tau_B \vdash c_i : \{\varepsilon\}\tau$.Since we can make y_i a fresh variable, we have $\Gamma, y_i : \tau_B, x_j : \tau' \vdash c_i : \{\varepsilon\}\tau$.Then by IH we have $\Gamma \vdash \{e_j/x_j\}e_i : \tau_A$ and $\Gamma, y_i : \tau_B \vdash \{e_j/x_j\}c_i : \{\varepsilon\}\tau$.By (T-Op) we have $\Gamma \vdash op(\{e_j/x_j\}e_i; y_i.\{e_j/x_j\}c_i) : \{\varepsilon\}\tau$ Therefore we have $\Gamma \vdash \{e_j/x_j\}(op(e_i; y_i.c_i)) : \{\varepsilon\}\tau$

(T-Seq)

$$\frac{\Gamma, x_j : \tau'' \vdash c_i : \{\varepsilon\}\tau \quad \Gamma, x_j : \tau'', x_i : \tau \vdash c'_i : \{\varepsilon\}\tau'}{\Gamma, x_j : \tau'' \vdash \text{do } x_i \leftarrow c_i \text{ in } c'_i : \{\varepsilon\}\tau'} \text{ (T-SEQ)}$$

By IH, we have $\Gamma \vdash \{e_j/x_j\}c_i : \{\varepsilon\}\tau$ Since we can choose x_i as a fresh variable, we have $\Gamma, x_i : \tau, x_j : \tau'' \vdash c'_i : \{\varepsilon\}\tau'$ Then by IH we have $\Gamma, x_i : \tau \vdash \{e_j/x_j\}c'_i : \{\varepsilon\}\tau'$

Then the result follows by (T-Seq)

(T-App) Follows directly by applying IH.

(T-Handle)

$$\frac{\begin{array}{l} h_i = \{ \text{return } x \mapsto c^r, op^1(x; k) \mapsto c^1, \dots, op^n(x; k) \mapsto c^n \} \\ \Gamma, x_j : \tau', x : \tau_A \vdash c^r : \{\varepsilon'\}\tau_B \\ \{\Sigma(op^i) = \tau_i \rightarrow \tau'_i \quad \Gamma, x_j : \tau', x : \tau_i, k : \tau'_i \rightarrow \{\varepsilon'\}\tau_B \vdash c^i : \{\varepsilon'\}\tau_B\}_{1 \leq i \leq n} \\ \Gamma, x_j : \tau' \vdash c_i : \{\varepsilon\}\tau_A \quad \varepsilon \setminus \{op^i\}_{1 \leq i \leq n} \subseteq \varepsilon' \end{array}}{\Gamma, x_j : \tau' \vdash \text{with } h_i \text{ handle } c_i : \{\varepsilon'\}\tau_B} \text{ (T-HANDLE)}$$

Then handling clauses bind variables x and k in the handling computation c^i , so we can make them fresh variables that do not appear in context Γ . Then we can apply IH to typing judgements in the premise.

(T-Embed) Follows by applying IH

(T-EmbedOp) The proof is similar to the case for (T-Op)

□

Lemma 7. *If $\Gamma \vdash c_i : \{\varepsilon\}\tau$ then $\overline{\Delta_i}(\varepsilon) = \varepsilon$*

Proof. By induction on derivation of $\Gamma \vdash c : \sigma$. (T-Ret) has a premise that ensures the lemma is correct. For other rules, the result is immediate by applying IH. \square

Lemma 8. *If $\varepsilon \leq_l \varepsilon'$, then $\varepsilon \setminus op \leq_l \varepsilon' \setminus op$*

Proof. By induction on $\varepsilon \leq_l \varepsilon$. The proof is straightforward. \square

Lemma 9. *If $op \leq_l \varepsilon$, then $op \leq_l \varepsilon \setminus op'$*

Proof. By induction on the derivation of $\varepsilon \leq_l \varepsilon$. If (R-Eff1) is used, then the proof is straightforward because the subset relation on the premise still holds. If (R-Eff2) is used, by inversion on (R-Eff2), we have $op \leq_l \varepsilon'$ and $\varepsilon' \leq_{l'} \varepsilon$. By IH we have $op \leq_l \varepsilon' \setminus op'$. By lemma 8 we have $\varepsilon' \setminus op' \leq_{l'} \varepsilon \setminus op'$. Then the result follows by (R-Eff2) \square

Lemma 10. *If $\tau \leq_l \tau'$ then $\tau' \leq_{rev(l)} \tau$*

Proof. By induction on the type relation rules. The proof is simple arguments that follow directly from IH. \square

Lemma 11. *(Preservation for expressions)*

For all agent i , If $\Gamma \vdash e_i : \tau$ and $e_i \mapsto e'_i$, then $\Gamma \vdash e'_i : \tau$.

Proof. By induction on derivation of $e_i \mapsto e'_i$

(E-Congruence) By inversion on the typing rule for embedded expressions, we have $\Gamma \vdash e_j : \tau'$. By IH, we have $\Gamma \vdash e'_j : \tau'$. Then we use (E-Contruence) to derive $\Gamma \vdash [e'_j]_j^\tau : \tau$

(E-Unit) Follows immediately from (T-Unit)

(E-Lambda) By inversion on (T-Embed), we have $\Gamma \vdash \lambda x_j : \tau'. c_j : \tau' \rightarrow \sigma'$, where $\tau' \rightarrow \sigma' \leq_{ji} \tau \rightarrow \sigma$.

By inversion on (R-Arrow), we have $\tau' \leq_{ji} \tau$ and $\sigma' \leq_{ji} \sigma$

By inversion on (T-Lambda), we have $\Gamma, x_j : \tau' \vdash c_j : \sigma'$. And since x_i is a fresh variable in c_j , we have $\Gamma, x_i : \tau, x_j : \tau' \vdash c_j : \sigma'$

By lemma 10, we have $\tau \leq_{ij} \tau'$, and therefore $\Gamma, x_i : \tau \vdash [x_i]_i^{\tau'} : \tau'$

Then we can use the substitution lemma to derive $\Gamma, x_i : \tau \vdash \{[x_i]_i^{\tau'}/x_j\}c_j : \sigma'$.

Then by (T-Embed), we have $\Gamma, x_i : \tau \vdash \{[x_i]_i^{\tau'}/x_j\}c_j]_j^\sigma : \sigma$

Then the result follows by (T-Lambda). \square

Lemma 12. *(Preservation for computations)*

If $\Gamma \vdash c_i : \{\varepsilon\}\tau$ and $c_i \longrightarrow c'_i$, then $\Gamma \vdash c'_i : \{\varepsilon\}\tau$

Proof. (Sketch) By induction on the derivation that $c_i \longrightarrow c'_i$. We proceed by the cases on the last step of the derivation.

1. E-Ret: By inversion, $\Gamma \vdash e_i : \tau$. By preservation of expressions and IH, we have $\Gamma \vdash e'_i : \tau$. Then we can use E-Ret to derive $\Gamma \vdash c'_i : \{\varepsilon\}\tau$
2. E-Op: Follow immediately from inversion and IH

3. E-EmbedOp1: Follow immediately from inversion and IH
4. E-EmbedOp2:

$$\frac{\overline{\Delta}_i(\varepsilon) = \varepsilon''}{[op_j]_l^\varepsilon(v_i; y_i.c_i) \longrightarrow [op_j]_l^{\varepsilon''}(v_i; y_i.c_i)} \text{ (E-EMBEDOP2)}$$

We have the typing rule as follows:

$$\frac{\Sigma(op) = \tau_A \rightarrow \tau_B \quad \Gamma \vdash e_i : \tau_A \quad \Gamma.y_i : \tau_B \vdash c_i : \{\varepsilon'\}\tau \quad \overline{\Delta}_i(\varepsilon) \subseteq \overline{\Delta}_i(\varepsilon') \quad \Gamma \vdash op \leq_{li} \varepsilon}{\Gamma \vdash [op]_l^\varepsilon(e_i; y_i.c_i) : \{\overline{\Delta}_i(\varepsilon')\}\tau} \text{ (T-EMBEDOP)}$$

Since $\overline{\Delta}_i(\varepsilon'') = \varepsilon''$ and $\varepsilon'' = \overline{\Delta}_i(\varepsilon)$, we have $\overline{\Delta}_i(\varepsilon'') \subseteq \overline{\Delta}_i(\varepsilon')$. Then we can use T-EmbedOp to derive $\Gamma \vdash [op]_l^{\varepsilon''}(e_i; y_i.c_i) : \{\overline{\Delta}_i(\varepsilon')\}\tau$

5. E-EmbedOp3: We have the typing rule as follows:

$$\frac{\Sigma(op) = \tau_A \rightarrow \tau_B \quad \Gamma \vdash e_i : \tau_A \quad \Gamma.y_i : \tau_B \vdash c_i : \{\varepsilon'\}\tau \quad \overline{\Delta}_i(\varepsilon) \subseteq \overline{\Delta}_i(\varepsilon') \quad \Gamma \vdash op \leq_{li} \varepsilon}{\Gamma \vdash [op]_l^\varepsilon(e_i; y_i.c_i) : \{\overline{\Delta}_i(\varepsilon')\}\tau} \text{ (T-EMBEDOP)}$$

By E-EmbedOp3, $op \in \overline{\Delta}_i(\varepsilon)$. So $op \in \overline{\Delta}_i(\varepsilon')$. By inversion on the typing rule, we have $\Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon'\}\tau$. By lemma 7, we have $\Gamma, y_i : \tau_B \vdash c_i : \{\overline{\Delta}_i(\varepsilon')\}\tau$. Then we can use T-Op to derive the designed result $\Gamma \vdash op(e_i; y_i.c_i) : \{\overline{\Delta}_i(\varepsilon')\}\tau$

6. E-EmbedOp4: We have the typing rule as follows:

$$\frac{\Sigma(op) = \tau_A \rightarrow \tau_B \quad \Gamma \vdash e_i : \tau_A \quad \Gamma.y_i : \tau_B \vdash c_i : \{\varepsilon'\}\tau \quad \overline{\Delta}_i(\varepsilon) \subseteq \overline{\Delta}_i(\varepsilon') \quad \Gamma \vdash op \leq_{li} \varepsilon}{\Gamma \vdash [op]_l^\varepsilon(e_i; y_i.c_i) : \{\overline{\Delta}_i(\varepsilon')\}\tau} \text{ (T-EMBEDOP)}$$

By lemma 9, we have $op \leq_{li} \varepsilon \setminus op'$. By inversion on the typing rule, we have $\Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon'\}\tau$ and $\varepsilon \subseteq \overline{\Delta}_i(\varepsilon')$. So $\varepsilon \setminus op' \subseteq \overline{\Delta}_i(\varepsilon')$. By lemma 7, we have $\Gamma, y_i : \tau_B \vdash c_i : \{\overline{\Delta}_i(\varepsilon')\}\tau$. Then we can apply T-EmbedOp again to derive $\Gamma \vdash [op]_l^{\varepsilon \setminus op'}(e_i; y_i.c_i) : \{\overline{\Delta}_i(\varepsilon')\}\tau$

7. E-App1: Follows immediately by T-App
8. E-App2: Follows immediately by T-App
9. E-App3: By inversion of T-App, we $\Gamma \vdash (\lambda x_i : \tau. c_i) : \tau \rightarrow \sigma, \Gamma \vdash v_i : \tau$. By inversion of T-Lam, $\Gamma, x_i : \tau \vdash c_i : \sigma$. By substitution lemma, we have $\Gamma \vdash \{v_i/x_i\}c_i : \sigma$.
10. E-Seq1: Follows immediately by T-Seq and IH.
11. E-Seq2: By inversion on T-Seq, we have $\Gamma \vdash \text{return } v_i : \{\varepsilon\}\tau$ and $\Gamma, x_i : \tau \vdash c'_i : \{\varepsilon\}\tau'$. By inversion on T-Ret, we have $\Gamma \vdash v_i : \tau$. Then by substitution lemma we have $\Gamma \vdash \{v_i/x\}c'_i : \{\varepsilon\}\tau'$.
12. E-Seq3:

$$\frac{}{\text{do } x \leftarrow op_i(v_i; y_i.c_i) \text{ in } c'_i \longrightarrow op_i(v_i; y_i. \text{do } x \leftarrow c_i \text{ in } c'_i)} \text{ (E-SEQ3)}$$

By inversion of T-Seq, we have $\Gamma \vdash op_i(v_i; y_i.c_i) : \{\varepsilon\}\tau$ and $\Gamma, x : \tau \vdash c'_i : \{\varepsilon\}\tau'$. By inversion on T-OP, we have $\Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon\}\tau$ and $op \in \varepsilon$ and $\Gamma \vdash v_i : \tau_A$. Then by T-Seq, we have $\Gamma, y_i : \tau_B \vdash \text{do } x \leftarrow c_i \text{ in } c'_i : \{\varepsilon\}\tau'$. Then we can use T-Op to derive $\Gamma \vdash op_i(v_i; y_i. \text{do } x \leftarrow c_i \text{ in } c'_i) : \{\varepsilon\}\tau'$.

13. E-Seq4

$$\frac{\Delta_i(\varepsilon) = \varepsilon \quad op \notin \varepsilon}{do\ x \leftarrow [op]_l^\varepsilon(v_i; y_i.c_i) \text{ in } c'_i \longrightarrow [op]_l^\varepsilon(v_i; y_i. do\ x \leftarrow c_i \text{ in } c'_i)} \text{ (E-SEQ4)}$$

$$\frac{\Gamma \vdash c_i : \{\varepsilon'\}\tau \quad \Gamma, x_i : \tau \vdash c'_i : \{\varepsilon'\}\tau'}{\Gamma \vdash do\ x_i \leftarrow c_i \text{ in } c'_i : \{\varepsilon'\}\tau'} \text{ (T-SEQ)}$$

By inversion on T-Seq, we have $\Gamma \vdash [op]_l^\varepsilon(v_i; y_i.c_i) : \{\varepsilon'\}\tau$ and $\Gamma, x : \tau \vdash c'_i : \{\varepsilon'\}\tau'$. Then by inversion on T-EmbedOp, we have $\Gamma, y_i : \tau_B \vdash c_i : \{\varepsilon'\}\tau$, $\overline{\Delta_i}(\varepsilon) \subseteq \varepsilon'$. Then by T-Seq, we have $\Gamma, y_i : \tau_B \vdash do\ x \leftarrow c_i \text{ in } c'_i : \{\varepsilon'\}\tau'$. Then by T-EmbedOp, we have $\Gamma \vdash [op]_l^\varepsilon(v_i; y_i. do\ x \leftarrow c_i \text{ in } c'_i) : \{\varepsilon'\}\tau'$

14. E-Handle1: Follows immediately by inversion on T-Handle and IH

15. E-Handle2: By T-Handle, we have $\Gamma \vdash$ with h_i handle return $v_i : \{\varepsilon'\}\tau_B$. By inversion on T-Handle, we have $\Gamma, x_i : \tau_A \vdash c'_i : \{\varepsilon'\}\tau_B$, and $\Gamma \vdash$ return $v_i : \{\varepsilon'\}\tau_A$. By inversion on T-Ret, we have $\Gamma \vdash v_i : \tau_A$. Then by substitution lemma, we have $\Gamma \vdash \{v_i/x_i\}c'_i : \{\varepsilon'\}\tau_B$.

16. E-Handle3

$$\frac{op(x_i; k) \mapsto c'_i \in h_i \quad \Sigma(op) = \tau_i \rightarrow \tau'_i}{\text{with } h_i \text{ handle } op(v; y_i.c_i) \longrightarrow \{v_i/x_i\}(\{\lambda y_i : \tau'_i. \text{with } h_i \text{ handle } c_i\}/k)\}c'_i} \text{ (E-HANDLE3)}$$

By T-Handle, we have $\Gamma \vdash$ with h_i handle $op(v; y_i.c_i) : \{\varepsilon'\}\tau_B$. By inversion on T-Handle, we have $\Gamma, x_i : \tau_i, k : \tau'_i \rightarrow \{\varepsilon'\}\tau_B \vdash c'_i : \{\varepsilon'\}\tau_B$, and $\Gamma \vdash op(v; y_i.c_i) : \{\varepsilon'\}\tau_A$. By inversion on T-Op, we have $\Gamma \vdash v_i : \tau_i$ and $\Gamma, y_i : \tau'_i \vdash c_i : \{\varepsilon'\}\tau_A$. By T-Handle, we have $\Gamma, y_i : \tau'_i \vdash$ with h_i handle $c_i : \{\varepsilon'\}\tau_B$. Then by T-Lam, we have $\Gamma \vdash \lambda y_i : \tau'_i. \text{with } h_i \text{ handle } c_i : \tau'_i \rightarrow \{\varepsilon'\}\tau_B$. Then, by substitution lemma, we have $\Gamma \vdash \{v_i/x_i\}(\{\lambda y_i : \tau'_i. \text{with } h_i \text{ handle } c_i\}/k)\}c'_i : \{\varepsilon'\}\tau_B$.

17. E-Handle4:

$$\frac{\Delta_i(\varepsilon) = \varepsilon \quad op \notin \varepsilon}{\text{with } h_i \text{ handle } [op]_l^\varepsilon(v_i, y_i.c_i) \longrightarrow [op]_l^\varepsilon(v_i; y_i. \text{with } h_i \text{ handle } c_i)} \text{ (E-HANDLE4)}$$

$$\frac{\begin{array}{l} h_i = \{ \text{return } x \mapsto c^r, op^1(x; k) \mapsto c^1, \dots, op^n(x; k) \mapsto c^n \} \\ \Gamma, x : \tau_A \vdash c^r : \{\varepsilon'\}\tau_B \\ \{\Sigma(op^i) = \tau_i \rightarrow \tau'_i \mid \Gamma, x : \tau_i, k : \tau'_i \rightarrow \{\varepsilon'\}\tau_B \vdash c^i : \{\varepsilon'\}\tau_B\}_{1 \leq i \leq n} \\ \Gamma \vdash c_i : \{\varepsilon''\}\tau_A \quad \varepsilon'' \setminus \{op^i\}_{1 \leq i \leq n} \subseteq \varepsilon' \end{array}}{\Gamma \vdash \text{with } h_i \text{ handle } c_i : \{\varepsilon'\}\tau_B} \text{ (T-HANDLE)}$$

By T-Handle, we have $\Gamma \vdash$ with h_i handle $[op]_l^\varepsilon(v; y_i.c_i) : \{\varepsilon'\}\tau_B$. By inversion on T-Handle, we have $\Gamma \vdash [op]_l^\varepsilon(v; y_i.c_i) : \{\varepsilon''\}\tau_A$ and $\varepsilon'' \setminus \{op^i\} \subseteq \varepsilon'$. By inversion on T-EmbedOp, we have $\Gamma \vdash v_i : \tau_i$, $\Gamma, y_i : \tau'_i \vdash c_i : \{\varepsilon''\}\tau_A$ and $\varepsilon \subseteq \varepsilon''$. Since ε doesn't contain any concrete operation, we have $\varepsilon \subseteq \varepsilon'' \setminus \{op^i\} \subseteq \varepsilon'$. Then by T-Handle, we have $\Gamma, y_i : \tau'_i \vdash$ with h_i handle $c_i : \{\varepsilon'\}\tau_B$. Then, we use T-EmbedOp to derive $\Gamma \vdash [op]_l^\varepsilon(v_i; y_i. \text{with } h_i \text{ handle } c_i) : \{\varepsilon'\}\tau_B$

18. E-Embed1: Follows immediately from Inversion and IH
19. E-Embed2: By typing rule, we have $\Gamma \vdash [\text{return } v_j]_l^{\{\varepsilon\}\tau} : \{\varepsilon\}\tau$. By inversion on the typing rule, we have $\Gamma \vdash \text{return } v_j : \{\varepsilon'\}\tau'$ such that $\{\varepsilon'\}\tau' \leq_{li} \{\varepsilon\}\tau$. By inversion on R-Sigma, we have $\tau' \leq_{li} \tau$. Then by T-EmbedExp, we have $\Gamma \vdash [v_j]_l^\tau : \tau$. Then by T-Ret, we have $\Gamma \vdash \text{return } [v_j]_l^\tau : \{\varepsilon\}\tau$. $\Gamma \vdash [\text{return } v_j]_l^{\{\varepsilon\}\tau} : \{\varepsilon\}\tau$
20. E-Embed3:

$$\frac{\Sigma(op) = \tau_A \rightarrow \tau_B}{[op(v_j; y_j.c_j)]_l^{\{\varepsilon\}\tau} \longrightarrow [op]_l^\varepsilon([v_j]_j^{\tau_A}; y_i.\{[y_i]_i^{\tau_B}/y_j\}[c_j]_l^{\{\varepsilon\}\tau})} \text{ (E-EMBED3)}$$

By typing rule, we have $\Gamma \vdash op(v_j; y_j.c_j) : \{\varepsilon'\}\tau'$, where $\{\varepsilon'\}\tau' \leq_{li} \{\varepsilon\}\tau$. By inversion on T-Op, we have $\Gamma \vdash v_j : \tau_A$, and $\Gamma, y_j : \tau_B \vdash c_j : \{\varepsilon'\}\tau'$. Then, by T-EmbedExp, we have $\Gamma \vdash [v_j]_j^{\tau_A} : \tau_A$. By substitution lemma, we have $\Gamma, y_i : \tau_B \vdash \{[y_i]_i^{\tau_B}/y_j\}c_j : \{\varepsilon'\}\tau'$. By T-Embed, we have $\Gamma, y_i : \tau_B \vdash \{[y_i]_i^{\tau_B}/y_j\}[c_j]_l^{\{\varepsilon\}\tau} : \{\varepsilon\}\tau$. Then we can use T-EmbedOp to derive $\Gamma \vdash [op]_l^\varepsilon([v_j]_j^{\tau_A}; y_i.\{[y_i]_i^{\tau_B}/y_j\}[c_j]_l^{\{\varepsilon\}\tau}) : \{\varepsilon\}\tau$.

21. E-Embed4:

$$\frac{\Sigma(op_k) = \tau_A \rightarrow \tau_B \quad \Delta_j(\varepsilon') = \varepsilon' \quad op \notin \varepsilon'}{[[op_k]_{l'}^{\varepsilon'}(v_j; y_j.c_j)]_l^{\{\varepsilon\}\tau} \longrightarrow [op_k]_{l',j,l}^\varepsilon([v_j]_j^{\tau_A}; y_i.\{[y_i]_i^{\tau_B}/y_j\}[c_j]_l^{\{\varepsilon\}\tau})} \text{ (E-EMBED4)}$$

By typing rule, we have $\Gamma \vdash [op_K]_{l'}^{\varepsilon'}(v_j; y_j.c_j) : \{\varepsilon''\}\tau''$, where $\{\varepsilon''\}\tau'' \leq_{li} \{\varepsilon\}\tau$. By inversion on T-EmbedOp, we have $\Gamma \vdash v_j : \tau_A$ and $\Gamma, y_j : \tau_B \vdash c_j : \{\varepsilon''\}\tau''$. Then, by T-EmbedExp, we have $\Gamma \vdash [v_j]_j^{\tau_A} : \tau_A$. By substitution lemma, we have $\Gamma, y_i : \tau_B \vdash \{[y_i]_i^{\tau_B}/y_j\}c_j : \{\varepsilon''\}\tau''$. By T-Embed, we have $\Gamma, y_i : \tau_B \vdash \{[y_i]_i^{\tau_B}/y_j\}[c_j]_l^{\{\varepsilon\}\tau} : \{\varepsilon\}\tau$. Then we use T-EmbedOp to derive $\Gamma \vdash [op_k]_{l',j,l}^\varepsilon([v_j]_j^{\tau_A}; y_i.\{[y_i]_i^{\tau_B}/y_j\}[c_j]_l^{\{\varepsilon\}\tau}) : \{\varepsilon\}\tau$.

□

Lemma 13. (Progress)

If $\emptyset \vdash c_i : \{\varepsilon\}\tau$ then either

1. $c_i \longrightarrow c'_i$
2. $c_i = \text{return } v_i$
3. $c_i = op(v_i; y_i.c'_i)$
4. $c_i = [op]_l^\varepsilon(v_i; y_i.c'_i)$

TODO: Proof

Definition 4.5.1. A i -computation c is oblivious to effect label f if $f \notin \text{Dom}(\delta_i)$, and for all subexpression $[e]_j^\tau$ and subcomputation $[c]_j^\sigma$, $f \notin \text{Dom}(\delta_j)$

Theorem 14. Let c_1 and c_2 be computations that are oblivious to the effect f . If $c_1 \approx c_2$, $c_1 \rightarrow c'_1$, $c_2 \rightarrow c'_2$, then $c'_1 \approx c'_2$. Furthermore, If e_1, e_2 oblivious to f , $e_1 \approx e_2$, $e_1 \rightarrow e'_1$, $e_2 \rightarrow e'_2$, then $e'_1 \approx e'_2$

The relation \approx is defined as follows:

$e \approx e$

$$\frac{}{x \approx x} \text{ (R-VAR)} \quad \frac{}{() \approx ()} \text{ (R-UNIT)}$$

$$\frac{c \approx c'}{\lambda x : \tau. c \approx \lambda x : \tau. c'} \text{ (R-LAM)} \quad \frac{e \approx e'}{[e]_l^\tau \approx [e']_l^\tau} \text{ (R-EMBEDEXP)}$$

 $c \approx c$

$$\frac{e \approx e'}{\text{return } e \approx \text{return } e'} \text{ (R-RET)} \quad \frac{e \approx e' \quad c \approx c'}{op''(e; y.c) \approx op''(e'; y.c')} \text{ (R-OP)}$$

$$\frac{c \approx c' \quad d \approx d'}{\text{do } x \leftarrow c \text{ in } d \approx \text{do } x \leftarrow c' \text{ in } d'} \text{ (R-SEQ)} \quad \frac{e_1 \approx e'_1 \quad e_2 \approx e'_2}{e_1 e_2 \approx e'_1 e'_2} \text{ (R-APP)}$$

$$\frac{h \approx h' \quad c \approx c'}{\text{with } h \text{ handle } c \approx \text{with } h' \text{ handle } c'} \text{ (R-HANDLE)} \quad \frac{c_j \approx c'_j}{[c_j]_l^\sigma \approx [c'_j]_l^\sigma} \text{ (R-EMBED)}$$

$$\frac{e \approx e' \quad c \approx c'}{[op'']_l^\varepsilon(e; y.c) \approx [op'']_l^\varepsilon(e'; y.c')} \text{ (R-EMBEDOP1)} \quad \frac{e \approx e' \quad c \approx c' \quad \exists i \in l, \delta_i(f) = op, op'}{[op]_l^\varepsilon(e; y.c) \approx [op']_l^\varepsilon(e'; y.c')} \text{ (R-EMBEDOP2)}$$

 $h \approx h$

$$\frac{c_r \approx c'_r \quad c_1 \approx c'_1, \dots, c_n \approx c'_n}{\{\text{return } x \mapsto c_r, op_1(x_1, k_1) \mapsto c_1, \dots, op_n(x_n, k_n) \mapsto c_n\} \approx \{\text{return } x \mapsto c'_r, op_1(x_1, k_1) \mapsto c'_1, \dots, op_n(x_n, k_n) \mapsto c'_n\}} \text{ (R-HANDLER)}$$

Figure 4.6: Definition of $\approx_{op, op'}$

Proof. (Sketch) By induction on derivation of $c_1 \approx c_2$ and $e_1 \approx e_2$

1. R-Ret: The only reduction rule that applies is E-Ret, so we have $e_1 \longrightarrow e'_1$ and $e_2 \longrightarrow e'_2$. By IH, we have $e'_1 \approx e'_2$. Then the result follows by R-Ret.
2. R-Op: The only reduction rule that applies is E-Op. The result is immediate by IH.
3. R-Seq: If the reduction rule is E-Seq1, then result is immediate by IH.
If the reduction rule is E-Seq2. Then we have $c_1 = \text{do } x \leftarrow \text{return } v_1 \text{ in } d_1, c_2 = \text{do } x \leftarrow \text{return } v_2 \text{ in } d_2$. By inversion, we have $v_1 \approx v_2$ and $d_1 \approx d_2$. So we have $\{v_1/x\}d_1 \approx \{v_2/x\}d_2$.
If the reduction rule is E-Seq3, then $c_1 = \text{do } x \leftarrow op(v_1; y.k_1) \text{ in } d_1, c_2 = \text{do } x \leftarrow op(v_2; y.k_2) \text{ in } d_2$. By inversion, we have $k_1 \approx k_2$ and $d_1 \approx d_2$. So $\text{do } x \leftarrow k_1 \text{ in } d_1 \approx \text{do } x \leftarrow k_2 \text{ in } d_2$. So $op(v_1; y. \text{do } x \leftarrow k_1 \text{ in } d_1) \approx op(v_2; y. \text{do } x \leftarrow k_2 \text{ in } d_2)$. The proof is similar for rule E-Seq4.
4. R-App: The cases for reduction rules E-App1 and E-App2 follows by IH. If reduction rule is E-App3. Then $c_1 = (\lambda x : \tau. d_1) v_1$ and $c_2 = (\lambda x : \tau. d_2) v_2$. By inversion we have $d_1 \approx d_2$ and $v_1 \approx v_2$. So we have $\{v_1/x\}d_1 \approx \{v_2/x\}d_2$.
5. R-Handle: If reduction rule is E-Handle1, then result follows by IH.
If the reduction rule is E-Handle2. Then $c_1 = \text{with } h_1 \text{ handle return } v_1$ and $c_2 = \text{with } h_2 \text{ handle return } v_2$. By inversion we have $h_1 \approx h_2, v_1 \approx v_2$.

Let $\text{return } c_{r1} \in h_1$ and $\text{return } c_{r2} \in h_2$. By inversion we have $c_{r1} \approx c_{r2}$. So $\{v_1/x\}c_{r1} \approx \{v_2/x\}c_{r2}$.

If the reduction rule is E-Handle3. Then $c_1 =$ with h_1 handle $op(v_1; y.k_1)$ and $c_2 =$ with h_2 handle $op(v_2; y.k_2)$. By inversion we have $v_1 \approx v_2$, $k_1 \approx k_2$ and $h_1 \approx h_2$. Then by equivalent rules we derive $c'_1 \approx c'_2$. The case for E-Handle4 is similar.

6. R-Embed: If reduction is E-Embed1, result is immediate by IH. If reduction rule is E-Embed2, then $c_1 = [\text{return } v_1]_l^{\{\varepsilon\}\tau}$ and $c_2 = [\text{return } v_2]_l^{\{\varepsilon\}\tau}$. It is easy to see $[v_1]^\tau \approx [v_2]^\tau$. So the result holds.

If the reduction rule is E-Embed3, Then $c_1 = [op(v_1; y.k_1)]\{\varepsilon\}\tau_l$ and $c_2 = [op(v_2; y.k_2)]\{\varepsilon\}\tau_l$. By inversion we have $v_1 \approx v_2$, $k_1 \approx k_2$. Then by equivalent rules we have $c'_1 \approx c'_2$. Same arguments apply for E-Embed4.

7. R-EmbedOp1: If reduction rule is E-EmbedOp1, then result follows by IH. If reduction rules is E-EmbedOp2 or E-EmbedOp3, reduction does not affect terms except effect annotation, so the equivalence relation still holds after reduction.
8. R-EmbedOp2: Reduction rules E-EmbedOp1 and E-EmbedOp2 are similar to the previous case. If the reduction rule is E-EmbedOp3, then by R-EmbedOp2, the operations op and op' are exported as effect f by some agent, and since current agent is oblivious to f , this case is impossible.

□

4.6 Abstraction Problem I

TODO: make the example readable

```

1
2
3 op : 1 -> 1
4
5 module b: B
6   f = op
7   def m() : {f} Unit
8     op ()
9   def handler(c: 1 -> {f} 1) : {} Int =
10     handle c () with
11       | op () -> 1
12       | return _ -> 0
13
14 // Example program
15 b.handler(
16   () => handle b.m() with
17     | op -> resume ()
18     | return _ -> ()
19 )
20
21 // Translation of Example Program
22 [\$\lambda$c: 1 -> {f} 1.
23   handle c() with
24     | op(x, k) -> return 1
25     | return x -> return 0)]$\^{(1 \rightarrow \{f\} 1) \rightarrow \{\} int}
26   _b$
27 ($\lambda$_:1.
28   handle [$\lambda$_:1. op()]\^{1 \rightarrow \{f\} 1}_b$ () with
29     | op(x, k) -> k ()
30     | return _ -> return ()
31
32 \\steps--->
33 ($\lambda$c: 1 -> {f} 1.
34   [handle [c]\^{1 \rightarrow \{f\} 1}_a$ () with
35     | op(x, k) -> return 1
36     | return x -> return 0)]$\^{ \{\} int}_b$)
37 ($\lambda$_:1.
38   handle [$\lambda$_:1. op()]\^{1 \rightarrow \{f\} 1}_b$ () with
39     | op(x, k) -> k ()
40     | return _ -> return ()
41
42 \\steps--->
43 [handle
44   [$\lambda$_:1.
45     handle [$\lambda$_:1. op()]\^{1 \rightarrow \{f\} 1}_b$ () with
46       | op(x, k) -> k ()
47       | return _ -> return () ]$\^{1 \rightarrow \{f\} 1}_a$ ()
48   with

```

```

48 | op(x, k) -> return 1
49 | return x -> return 0)]$^{ \{\} int}_b$
50
51 \\steps--->
52 [handle
53   $\lambda$.1.
54   [handle [$\lambda$.1. op()]]$^{1 \rightarrow \{f\} 1}_b$ () with
55     | op(x, k) -> k ()
56     | return _ -> return () ]$^{ \{f\} 1}_a$ ()
57   with
58     | op(x, k) -> return 1
59     | return x -> return 0)]$^{ \{\} int}_b$
60
61 \\steps--->
62 [handle
63   [handle [$\lambda$.1. op()]]$^{1 \rightarrow \{f\} 1}_b$ () with
64     | op(x, k) -> k ()
65     | return _ -> return () ]$^{ \{f\} 1}_a$
66   with
67     | op(x, k) -> return 1
68     | return x -> return 0)]$^{ \{\} int}_b$
69
70
71 \\steps--->
72 [handle
73   [handle $\lambda$.1. [op()]]$^{ \{f\} 1}_b$ () with
74     | op(x, k) -> k ()
75     | return _ -> return () ]$^{ \{f\} 1}_a$
76   with
77     | op(x, k) -> return 1
78     | return x -> return 0)]$^{ \{\} int}_b$
79
80 \\steps--->
81 [handle
82   [handle [op()]]$^{ \{f\} 1}_b$ with
83     | op(x, k) -> k ()
84     | return _ -> return () ]$^{ \{f\} 1}_a$
85   with
86     | op(x, k) -> return 1
87     | return x -> return 0)]$^{ \{\} int}_b$
88
89 \\steps--->
90 [handle
91   [handle op$^f_b$(), y.return y) with
92     | op(x, k) -> k ()
93     | return _ -> return () ]$^{ \{f\} 1}_a$
94   with
95     | op(x, k) -> return 1
96     | return x -> return 0)]$^{ \{\} int}_b$
97
98 \\steps--->
99 [handle

```



```

100     [op$^f_b$(), y. handle return y with
101         | op(x, k) -> k ()
102         | return _ -> return ()
103     )]$^{\{f\} 1}_a$
104     with
105     | op(x, k) -> return 1
106     | return x -> return 0)]$^{\{\} int}_b$
107
108 \steps--->
109 [handle
110     op$^f_{ba}$(), y. [handle return y with
111         | op(x, k) -> k ()
112         | return _ -> return ()]$^{\{f\} 1}_a$
113     )
114     with
115     | op(x, k) -> return 1
116     | return x -> return 0)]$^{\{\} int}_b$
117
118 \steps--->
119 [handle
120     op(), y. [handle return y with
121         | op(x, k) -> k ()
122         | return _ -> return ()]$^{\{f\} 1}_a$
123     )
124     with
125     | op(x, k) -> return 1
126     | return x -> return 0)]$^{\{\} int}_b$

```


Chapter 5

Algebraic Effects with Module Systems

5.1 Programming with Abstract Effects

In section 3.1.2, we defined a resource type `Var` that provides two public functions, `get` and `set`, that are annotated with effects, but the concrete implementation that causes the effects are not revealed to the client. This construct ensures that any program that accesses to a value of type `Var` is annotated with the correct effect. Therefore we can see that, in a restrictive effect system, it is useful to be able to hide the implementation of a computational effect.

In this section we will show that this form of implementation hiding is also important for building a modular software with algebraic effects and handlers through an example of mutable state. Mutable state as an algebraic effects is often represented by a `state` effect with two operations `get` and `set`. In this example we assume that our mutable state can store or access a `int`.

```
1 operation get : Unit -> Int
2 operation set : Int -> Unit
```

The handler of a state effect is usually defined in the following way:

```
1 handler hstate = {
2   return x -> λ_:Int. return x
3   get(_; k) -> λs:Int. (k s) s
4   set(s; k) -> λ_:Int. (k ()) s
5 }
```

This handler would transform the handled computation into a lambda expression that receives a state as an argument. In the return clause, the argument is ignored. In the clause handling `get`, the state argument is passed into the continuation k . Since the continuation $k\ s$ is already transformed into a function that expects a state, we pass s to the computation $k\ s$. The clause for the `set` operation is similar except that we ignore the state argument, but pass the argument obtained from the `set` operation.

Now we consider the module of a single variable introduced in section 3.1.2, where read and write to the variable cannot bypass the effect checking because the implementation details are hidden by the interface `Var`. If we would write a similar module in our core calculus, it is natural to extend our calculus with record type and universal types, and implement the module as following:

```

1 val var =
2   < read  $\hookrightarrow \lambda x:\text{Unit}. \text{get}(); y.\text{return } y$  ,
3     write  $\hookrightarrow \lambda x:\text{Int}. \text{set}(x; y.\text{return } ())$  ,
4     handle  $\hookrightarrow \lambda c:\text{Unit} \rightarrow \{\text{state}\} \text{Int}. (\text{with hstate handle } c ()) 0 >$ 

```

The module provides functions `read`, `write`, and a handler function `handle` that determines the semantics for operations in functions `read` and `write`. However, this encoding of the module `var` does not enforce that the operations `get` and `set` are always handled by the `handle` function. In fact, any client that calls function `read` and `write` can write its own handler to handle the effects. So in order to tackle with this issue, we use the existential type introduced in section 5.3 to define the module:

```

1 val var =
2   pack
3   <{get, set},
4     < read  $\hookrightarrow \lambda x:\text{Unit}. \text{get}(); y.\text{return } y$  ,
5       write  $\hookrightarrow \lambda x:\text{Int}. \text{set}(x; y.\text{return } ())$  ,
6       handle  $\hookrightarrow \lambda c:\text{Unit} \rightarrow \{\text{state}\} \text{Int}. (\text{with hstate handle } c ()) 0 >$ 
7   > as  $\exists \text{state}. (\text{read} : \text{Unit} \rightarrow \{\text{state}\} \text{Int}) * (\text{write} : \text{Int} \rightarrow \{\text{state}\} \text{Unit})$ 
   * (handle :  $\text{Unit} \rightarrow \{\text{state}\} \text{Int} \rightarrow \{\} \text{Int}$ )

```

This encoding of module defines an abstract effect `state` on top of the effects `get` and `set`, and export the module as an existential type that hides the definition of effect `state`, therefore enforcing that the client of this module can only handle effect `state` by calling the `handle` function.

The embedding design presented in section 4.2 helps to ensure that the abstraction does not break. Imagine a client of module `var` that calls `read` and handles it by calling `handle`.

```

1 open var as (state, v) in (v.handle v.read)

```

In this example, the `state` effect in function `read` is correctly handled by the handler `handle`, and the result of the computation is 0. Note that the abstraction barrier is still preserved if some handler attempts to handle the abstracted effect. For example, let handler `hget` be a handler that handles effect `get`. And the client code uses `hget` to handle the effect in method `read`.

```

1 handler hget = {
2   return x -> x
3   get (x; k) -> 1
4 }
5 open var as (state, v) in
6   (v.handle ( $\lambda x:\text{Unit}. \text{with hget handle } v.\text{read} ()$ ))

```

Assume that the `open` expression is an *i*-expression. When opening the module `var`, we assign the opened expression `v` to a new agent *j* that knows the definition of the `state` effect. Because the `hget` handler is evaluating in the agent *i*, it would not be able to handle the effect `state` of the function `v.read`. Instead, the `state` effect will be handled by the handler provided by the module `var`, therefore ensuring that the abstraction information is not leaked.

5.2 Future Work

5.2.1 Parametric Polymorphism

We introduced polymorphic effects in chapter 3 as a part of our restrictive effect system. However, we did not include polymorphism in our agent-based core calculus. Parametric polymorphism on effects would significantly increase the expressiveness of the language. For example, in the example of the `var` module presented in the previous section, the type of the argument to the `handle` function is `Unit -> {state} Int`. So it restricts the handled computation to only have the `state` effect, and is therefore unrealistic as we may want to handle computations with other effects as well.

```

1 val var =
2   pack
3   <{get, set},
4   < read ↦ λx:Unit. get(); y.return y),
5   write ↦ λx:Int. set(x; y.return ()),
6   handle ↦ λc: Unit → {state} Int. (with hstate handle c ()) 0 >
7 > as ∃state. (read : Unit -> {state} Int) * (write : Int -> {state} Unit)
   * (handle : Unit -> {state} Int -> {} Int)

```

Therefore we would desire to add universal quantifications for effect variables to the system. However, unlike existential quantification, which is a straightforward extension to our calculus, the universal effect introduces a problem that breaks the abstraction barrier. The following example illustrates the problem brought by parametric polymorphism on effects:

```

1 type B
2   effect E {
3     def op() : Unit
4   }
5
6 module b : B
7   ...
8
9 type A
10  effect E
11  def handle[F](c: Unit -> {F, this.E} Unit): {F} Unit
12  def m(): {this.E} Unit
13
14 module a: A
15  effect E = {b.E}
16  def handle[F](c: Unit -> {F, this.E} Unit): {F} Unit
17    handle
18      c()
19    with
20      b.op() -> resume ()
21  def m(): {this.E} Unit
22    b.op()
23

```

The module `a` defines a polymorphic `handle` function that handles a computation with effect `a.E` and a polymorphic effect `F`. Since effect `a.E` is defined abstract in the type `A`, the client

of this module should not observe that fact that effect $a.E$ is equivalent to $b.E$. However, the client1 passes effect $b.E$ as the polymorphic effect into the handle function, and because the implementation of handle function handles the effect $b.E$, client1 would be surprised by the fact that the effect $b.E$ is handled. In comparison, the client2 code passes an unrelated effect $c.E$ to the handling function, and observes that the operation of effect $c.E$ is not handled.

```

1 //Client1: Prints nothing
2 handle
3   a.handle[b.E] (
4     () => b.op(); a.m()
5   )
6 with
7   b.op() -> print "desired output"
8
9 //Client2: Prints "desired output"
10 handle
11   a.handle[c.E] (
12     () => c.op(); a.m()
13   )
14 with
15   c.op() -> print "desired output"

```

5.3 Existential Type for Algebraic Effect

In the previous sections, we have introduced the core calculus of abstract algebraic effects via embeddings. However it is impractical for requiring programmers to explicitly annotate each program components with embeddings. So we present a top level language where the annotations are implicitly added during evaluation of the program.

According to Mitchell and Plotkin [21], there is a correspondence between abstraction data types and existential types. Existential types are often used as a foundation for expressing type abstraction in module systems. The calculus introduced in this section contains a form of existential type that provide abstraction mechanisms for algebraic effects. The this form of expression of existential type would generate new agents during evaluation of program and automatically separate program components with different knowledge, so programmers would not need to explicitly work with agents and embeddings.

5.3.1 Syntax

Most of the syntax remains the same for our new language. The existential type $\exists f. \tau$ is added as an expression type. The intuition of the type is that the value of this type is a value of type $\{\varepsilon/f\}\tau$ for some effect type ε .

There are two new forms of expressions: The introduction form of the existential type, $\text{pack } (\varepsilon, e) \text{ as } \exists f. \tau$ and the elimination form, $\text{open } e \text{ as } (f, x) \text{ in } e'$. The `pack` expression creates existential package that contains an effect type ε and an expression e . The `open` expression opens up an existential package and substitute the expression in the package for the variable x .

(agents)	$i, j ::= \{1 \dots n\}$
(lists)	$l ::= i \mid il$
(expression types)	$\tau ::= 1 \mid \tau \rightarrow \sigma \mid \exists f. \tau$
(computation types)	$\sigma ::= \{\varepsilon\}\tau$
(effect types)	$\varepsilon ::= \cdot \mid f, \varepsilon \mid op, \varepsilon$
(i values)	$v_i ::= ()_i \mid \lambda x_i : \tau. c_i \mid \text{pack } (\varepsilon, v) \text{ as } \exists f. \tau$
(i expression)	$e_i ::= x_i \mid v_i \mid [e_j]_l^\tau \mid \text{pack } (\varepsilon, e) \text{ as } \exists f. \tau \mid \text{open } e \text{ as } (f, x) \text{ in } e'$
(i computation)	$c_i ::= \text{return } e_i \mid op(e_i, y. c_i) \mid \text{do } x \leftarrow c_i \text{ in } c'_i \mid e_i e'_i \mid \text{with } h_i \text{ handle } c_i$ $\mid [c_j]_l^\sigma \mid [op]_l^\varepsilon(e_i, y_i. c_i)$
(i handler)	$h_i ::= \text{handler } \{\text{return } x_i \mapsto c_i^r, op^1(x_i^1, k^1) \mapsto c_i^1 \dots op^n(x_i^n, k^n) \mapsto c_i^n\}$

Figure 5.1: Syntax for Existential Effects

$$\boxed{\langle \{\Delta\}, e \rangle \mapsto \langle \{\Delta\}, e \rangle}$$

$$\frac{\langle \{\Delta\}, e \rangle \mapsto \langle \{\Delta'\}, e' \rangle}{\langle \{\Delta\}, \text{pack } (\varepsilon, e) \text{ as } \exists f. \tau \rangle \mapsto \langle \{\Delta'\}, \text{pack } (\varepsilon, e') \text{ as } \exists f. \tau \rangle} \text{ (E-PACK)}$$

$$\frac{f \text{ fresh} \quad j \text{ fresh} \quad \Delta'_j = \Delta_i[f \mapsto \varepsilon] \quad \forall \Delta_k \in \{\Delta\}, \Delta'_k = \Delta_k[f \mapsto f]}{\langle \{\Delta\}, \text{open } (\text{pack } (\varepsilon, e) \text{ as } \exists f. \tau) \text{ as } (f, x) \text{ in } e' \rangle \mapsto \langle \{\Delta'\}, \{[e]_j^\tau / x\} e' \rangle} \text{ (E-OPEN)}$$

$$\frac{}{\langle \{\Delta\}, [\text{pack } (\varepsilon, v) \text{ as } \exists f. \tau]_j^{\exists f. \tau'} \rangle \mapsto \langle \{\Delta\}, \text{pack } (\varepsilon, [v]_j^{\{\varepsilon/f\}\tau'}) \text{ as } \exists f. \tau' \rangle} \text{ (E-EMBEDPACK)}$$

Figure 5.2: Additional Dynamic Semantics for Existential Type

We only introduce one form of value: the existential package $\text{pack } (\varepsilon, v) \text{ as } \exists f. \tau$

5.3.2 Dynamic Semantics

The semantics for existential type $\exists f. \tau$ hides the definition of the effect label f from the client of the value of this type. We leverage the our previous design of multi-agent calculus to achieve information hiding. However, the previous design assumes that the type information for each agents are predetermined and does not change during evaluation. Since existential types generate new abstraction boundaries, we need a different semantics that allow agents to be created during evaluation. Therefore, we modify the reduction rule to evaluate a pair that contains both the expression to evaluate and a context of type information. The idea to keep track of type information while evaluating terms was used by Grossman et al. [10] to encode parametric polymorphism in their system.

We introduce the notation $\{\Delta\}$ to express a list of type maps for all agents in the context

$$\boxed{\{\Delta\} \mid \Gamma \vdash e_i : \tau}$$

$$\frac{\{\Delta\} \mid \Gamma \vdash e : \{\varepsilon/f\}\tau}{\{\Delta\} \mid \Gamma \vdash \text{pack}(\varepsilon, e) \text{ as } \exists f. \tau : \exists f. \tau} \text{ (T-PACK)}$$

$$\frac{\{\Delta\} \mid \Gamma \vdash e : \exists f. \tau \quad \forall \Delta_i \in \{\Delta\}, \Delta'_i = \Delta_i[f \mapsto f] \quad \{\Delta'\} \mid \Gamma, x : \tau \vdash e' : \tau'}{\{\Delta\} \mid \Gamma \vdash \text{open } e \text{ as } (f, x) \text{ in } e' : \tau'} \text{ (T-OPEN)}$$

Figure 5.3: Additional Static Semantics for Existential Effects

$\{\Delta_1, \dots, \Delta_n\}$. The type information of each agents can change, and new agents can be generated, the evaluation judgment becomes $\langle \{\Delta\}, e \rangle \mapsto \langle \{\Delta\}, e \rangle$.

(E-Pack) shows the congruence rule for reduction of a pack expression. (E-Open) opens an existential package: This rule requires f to be a fresh label, which achievable by doing alpha conversion in e' . j is a fresh agent. A new type map for agent j extends the type map for i by mapping f to ε . Every existing type map in $\{\Delta\}$ is extended by mapping f to itself. Finally, $[e]_j^\tau$ is substituted for x in e' .

(E-EmbedPack) shows how the embedding interacts with existential packages. The existential package is lifted out of the embedding, and the value v becomes an embedded j -value with type annotation $\{\varepsilon/f\}\tau'$.

The reduction rules for remaining terms are not changed by the introduction of $\{\Delta\}$ and are therefore not shown.

5.3.3 Static Semantics

The typing rules also requires the set of type maps, so the judgment have the form $\{\Delta\} \mid \Gamma \vdash e : \tau$ (T-Pack) assigns the type $\exists f. \tau$ to the existential package if the expression e has type $\{\varepsilon/f\}\tau$. The rule (T-Open) assigns the type τ' to the open expression if e has the existential type $\exists f. \tau$ and e' has type τ' given that the context is extended with variable x , and the set of type maps is extended with the effect label f .

5.3.4 Type Relation

5.3.5 Type Safety

Lemma 15. (*Preservation for expression*)

If $\{\Delta\} \mid \Gamma \vdash e : \tau$ and $\langle \{\Delta\}, e \rangle \mapsto \langle \{\Delta'\}, e' \rangle$, then $\{\Delta'\} \mid \Gamma \vdash e' : \tau$

Proof. By rule induction on the dynamic semantics of expressions.

1. (E-Congruence): By inversion on typing judgement and applying IH.
2. (E-Unit): By directly applying (T-Unit)

3. (E-Lambda):

$$\frac{}{[\lambda x_j : \tau'. c_j]_j^{\tau \rightarrow \sigma} \longrightarrow \lambda x_i : \tau. [\{[x_i]_i^{\tau'} / x_j\} c_j]_{jl}^\sigma} \text{ (E-LAMBDA)}$$

By inversion on (T-EmbedExp), we have $\{\Delta\} \mid \Gamma \vdash [\lambda x_j : \tau'. c_j]_j^{\tau \rightarrow \sigma} : \tau \rightarrow \sigma$, and $\{\Delta\} \mid \Gamma \vdash \lambda x_j : \tau'. c_j : \tau' \rightarrow \sigma'$, where $\{\Delta\} \mid \Gamma \vdash \tau' \rightarrow \sigma' \leq_j \tau \rightarrow \sigma$. By inversion on (T-Lam), we have $\{\Delta\} \mid \Gamma, x_j : \tau' \vdash c_j : \sigma$. Then by substitution lemma, we have $\{\Delta\} \mid \Gamma, x_i : \tau \vdash \{[x_i]_i^{\tau'} / x_j\} c_j : \sigma'$. By (T-Embed), $\{\Delta\} \mid \Gamma, x_i : \tau \vdash [\{[x_i]_i^{\tau'} / x_j\} c_j : \sigma']_j^\sigma : \sigma$. Then the result follows by (T-Lam).

4. (E-Pack): By inversion and IH.

5. (E-Open): By inversion on (T-Pack), we have $\{\Delta\} \mid \Gamma \vdash e : \{\varepsilon/f\}\tau$. Then by (T-EmbedExp), we have $\{\Delta'\} \mid \Gamma \vdash [e]_j^\tau : \tau$. By inversion on (T-Open), $\{\Delta\} \mid \Gamma, x : \tau \vdash e' : \tau'$. Since j is fresh, e' doesn't contain any j term, so the type information from agent j doesn't affect the typing of e' . Therefore, $\{\Delta'\} \mid \Gamma, x : \tau \vdash e' : \tau'$. Finally, by substitution lemma, we have $\{\Delta'\} \mid \Gamma \vdash \{[e]_j^\tau / x\} e' : \tau'$.6. (E-EmbedPack): By inversion on (T-EmbedExp), we have $\{\Delta\} \mid \Gamma \vdash \text{pack}(\varepsilon, v) \text{ as } \exists f. \tau : \exists f. \tau$, and $\{\Delta\} \mid \Gamma \vdash \exists f. \tau \leq_j \exists f. \tau'$. By type relation, we have $\tau \leq_j \tau'$. Then by inversion on (T-Pack), we have $\{\Delta\} \mid \Gamma \vdash v : \{\varepsilon/f\}\tau$. Then by T-EmbedExp, we have $\{\Delta\} \mid \Gamma \vdash [v]_j^{\{\varepsilon/f\}\tau} : \{\varepsilon/f\}\tau'$. Then by (T-Pack), we get $\{\Delta\} \mid \Gamma \vdash \text{pack}(\varepsilon, [v]_j^{\{\varepsilon/f\}\tau'}) \text{ as } \exists f. \tau' : \exists f. \tau'$

□

Lemma 16. (Progress)

If $\{\Delta\} \mid \Gamma \vdash e : \tau$ then either e is a value or there exists e' and $\{\Delta'\}$ such that $\langle \{\Delta\}, e \rangle \mapsto \langle \{\Delta'\}, e' \rangle$

TODO: Proof

5.4 Wyvern with Algebraic Effects and Handlers

Finish this section

5.4.1 Syntax

In this section we extend the formalization of Wyvern programming language with abstract algebraic effects and handlers. In figure 5.4, we show an updated version of the Wyverns's object-oriented, and it is extended with several new constructs. The expression x and location l is the same from the original Wyvern language. The initialization $\text{new } (x \Rightarrow \bar{d})$ initializes an object with a list of declarations \bar{d} .

The embedding of expression $[e]_j^\tau$ is added to separate code with different knowledge of effect abstraction, as shown in chapter 4. We also introduce the type coercion $e \triangleright \tau$ to make the subtyping of our calculus explicit.

$e ::=$	x	$d ::=$	$m(x : \tau) = c$
	l		effect $E = \varepsilon$
	new $(x \Rightarrow \bar{d})$	$\varepsilon ::=$	$\cdot \mid e.E, \varepsilon \mid op, \varepsilon$
	$[e]_j^\tau$	$\tau ::=$	$\{x \Rightarrow \bar{\rho}\}$
	$e \triangleright \tau$	$\Gamma ::=$	$\emptyset \mid \Gamma, x : \tau$
$c ::=$	return e	$\rho ::=$	$m : \tau \rightarrow \sigma$
	$e.m(e)$		effect E
	$op(e; y.c)$		effect $E = \varepsilon$
	with h handle c	$\sigma ::=$	$\{\varepsilon\}\tau$
	$[c]^\sigma$		
	$[op]^\varepsilon(e; y.c)$		

Figure 5.4: Wyvern’s object-oriented core syntax with algebraic effects and handlers

There are two forms of declarations. The method declaration $m(x : \tau) = c$, and the effect declaration **effect** $E = \varepsilon$. The object type has the form $\{x \Rightarrow \rho\}$, where ρ is the declaration type. A declaration type can be an method type $m : \tau \rightarrow \sigma$, an abstract effect **effect** E , or a concrete effect **effect** $E = \varepsilon$.

5.4.2 Operational Semantics

5.4.3 Typing Rules

5.4.4 Type Safety

Lemma 17. (*Preservation for expressions*) *If $\{\Delta\} \mid \mu \mid \Gamma \vdash e : \tau$ and $\langle \{\Delta\} \mid \mu \mid e \rangle \mapsto \langle \{\Delta'\} \mid \mu' \mid e' \rangle$, then $\{\Delta'\} \mid \mu' \mid \Gamma \vdash e' : \tau$*

$$\boxed{\langle \{\Delta\} \mid \mu \mid e \rangle \mapsto \langle \{\Delta\}, \mu, e \rangle}$$

$$\frac{l \text{ fresh} \quad \forall \text{effect} \quad E = \varepsilon \in \bar{d} \quad \forall \Delta_k \in \{\Delta\} \quad \Delta'_k = \Delta_k[l.E \mapsto \varepsilon]}{\langle \{\Delta\} \mid \mu \mid \text{new}(x \Rightarrow \bar{d}) \mapsto \langle \{\Delta'\} \mid \mu, l \mapsto \{[l/x]\bar{d}\} \mid l \rangle} \quad (\text{E-NEW})$$

$$\frac{\langle \{\Delta\} \mid \mu \mid e \rangle \mapsto \langle \{\Delta'\} \mid \mu' \mid e' \rangle}{\langle \{\Delta\} \mid \mu \mid [e]_j^\tau \rangle \mapsto \langle \{\Delta'\} \mid \mu' \mid [e']_j^\tau \rangle} \quad (\text{E-EMBEDEXP1})$$

$$\frac{\begin{array}{l} \tau = \{x \Rightarrow \bar{\rho}\} \quad l \mapsto \{x \Rightarrow \bar{d}\} \in \mu \\ (\forall \rho_i \in \bar{\rho} \quad \rho_i = m : \tau' \rightarrow \sigma \quad m(y : \tau'') = c \in \bar{d} \quad d'_i = (m(z : \tau') = [\{[z]_i^{\tau''}/y\}c]_j^\sigma)) \\ (\forall \text{effect} \quad E = \varepsilon \in \bar{d} \quad \forall \Delta_k \in \{\Delta\} \quad \Delta'_k = \Delta_k[l'.E \mapsto l.E]) \end{array}}{\langle \{\Delta\} \mid \mu \mid [l]_j^\tau \rangle \mapsto \langle \{\Delta'\} \mid \mu, l' \mapsto \{x \Rightarrow \bar{d}'\} \mid l' \rangle} \quad (\text{E-EMBEDEXP2})$$

$$\frac{\langle \{\Delta\} \mid \mu \mid e \rangle \mapsto \langle \{\Delta'\} \mid \mu' \mid e' \rangle}{\langle \{\Delta\} \mid \mu \mid e : \tau \rangle \mapsto \langle \{\Delta'\} \mid \mu' \mid e' : \tau \rangle} \quad (\text{E-COERCION1})$$

$$\frac{j \text{ fresh} \quad (\forall \text{effect} \quad E \in \tau \quad (\forall \Delta_k \in \{\Delta\} \quad \Delta_k(l.E) = l.E) \quad \Delta_j(l.E) = \varepsilon)}{\langle \{\Delta\} \mid \mu \mid l : \tau \rangle \mapsto \langle \{\Delta'\} \mid \mu' \mid [l]_j^\tau \rangle} \quad (\text{E-COERCION2})$$

Figure 5.5: Dynamic semantics for Wyvern extended with algebraic effects

$$\boxed{\{\Delta\} \mid \Gamma \vdash e_i : \tau}$$

$$\overline{\{\Delta\} \mid \mu \mid \Gamma \vdash x : \Gamma(x)} \quad (\text{T-VAR})$$

$$\frac{\{\Delta\} \mid \mu \mid \Gamma \vdash \{\bar{d}\} : \tau}{\{\Delta\} \mid \mu, l \mapsto \{\bar{d}\} \mid \Gamma \vdash l : \tau} \quad (\text{T-LOC})$$

$$\frac{\{\Delta\} \mid \mu \mid \Gamma \vdash d_i : \rho_i}{\{\Delta\} \mid \mu \mid \Gamma \vdash \text{new}(\bar{d}) : \bar{\rho}} \quad (\text{T-NEW})$$

$$\frac{\{\Delta\} \mid \mu \mid \Gamma \vdash e : \tau' \quad \tau' \leq_{ji} \tau}{\{\Delta\} \mid \mu \mid \Gamma \vdash [e]_j^\tau : \tau} \quad (\text{T-EMBEDEXP})$$

$$\frac{\{\Delta\} \mid \mu \mid \Gamma \vdash e : \tau' \quad \tau' <: \tau}{\{\Delta\} \mid \mu \mid \Gamma \vdash e \triangleright \tau : \tau} \quad (\text{T-COERCION})$$

Figure 5.6: Typing rules for Wyvern extended with algebraic effects

$$\begin{array}{c}
 \boxed{\tau <: \tau} \\
 \\
 \frac{\forall i, \rho_i <: \rho'_i}{\{\bar{\rho}\} <: \{\rho'\}} \text{ (S-DEPTH)} \\
 \\
 \boxed{\rho <: \rho} \\
 \\
 \frac{}{\rho <: \rho} \text{ (S-REFL)} \\
 \\
 \frac{}{\text{effect } E = \varepsilon <: \text{effect } E} \text{ (S-EFFECT)}
 \end{array}$$

Figure 5.7: Subtyping rules for Wyvern extended with algebraic effects

Chapter 6

Related Work and Conclusion

6.1 Related Work

6.1.1 Restrictive Effect Systems

A restrictive effect system considers effects that are built into the language, such as reading and writing states or exceptions, and provide a way to restrict the usage of such effects. Our effect system introduced in chapter 3 is a restrictive effect system. Restrictive effect systems were first proposed by Lucassen [17] to track reads and writes to memory. Then Lucassen and Gifford [18] extended this effect system to support polymorphism. Restrictive Effects have since been used for a wide variety of purposes, including exceptions in Java [11] and asynchronous event handling [5].

6.1.2 Bounded Effect Polymorphism.

A limited form of bounded effect polymorphism were explored by Trifonov and Shao [29], who bound effect parameters by the resources they may act on; however, the bound cannot be another arbitrary effect, as in our system. Long et al. [16] use a form of bounded effect polymorphism internally but do not expose it to users of their system.

6.1.3 Subeffecting.

Some effect systems, such as Koka [12], provide a built-in set of effects with fixed sub-effecting relationships between them. Rytz et al. [27] supports more flexibility via an extensible framework for effects. Users can plug in their own domain of effects, specifying an effect lattice representing sueffecting relationships. Each plugin is monolithic. In contrast, our effect members allow new effects to be incrementally added and related to existing effects using declared subeffect bounds.

6.1.4 Path-dependent Effects

The Effekt Library BRACHTHÄUSER et al. [4] explores algebraic effects as a library of the Scala programming language. Since their effect system is built on the path-dependent type sys-

tem of Scala, it bears some similarities to our system. However, their work is largely orthogonal to our contributions due to following reasons:

Our goal is to check the effects of general purpose code. In contrast, Brachthäuser’s approach requires all effect-checked code to be written in a monad. This is required to support control effects (i.e. prescriptive effects), but it is not an incidental difference: it is also an integral part of their static effect checking system, because monads are the way that they couple Scala’s type members (which provide abstraction and polymorphism) to effects. Their paper therefore, does not solve the problem of soundly checking effects for non-monadic code. Most code in Scala and Wyvern—let alone more conventional languages such as Java—is non-monadic, for good reasons: monads are restrictive and, for some kinds of programming, quite awkward. Programmers may be willing to use monads narrowly to get the benefits of control effects that Brachthäuser et al. support, but outside the Haskell community, it does not seem likely they would be willing to use them at a much broader scale for the purpose of descriptive effects.

Our approach provides abstraction and polymorphism for descriptive effects. As discussed above, Brachthäuser et al.’s leverage of Scala’s type members provides abstraction and polymorphism only for prescriptive effects, and only in the context of monadic code. Even setting aside the issue of monads, above, it is unclear how their approach can provide abstraction for descriptive effects. The reason is that their abstraction works backwards from the kind we need. For example, we want to be able to implement a logger in terms of file I/O—and hide the fact that it is implemented that way. The natural way to start would be to model file I/O in their system as a set of effect operations that “handle” I/O operations at the top level (their system does not provide support for this, so it would have to be added). A logger library could then provide a set of “log” effects and a handler for them, implemented in terms of the top-level I/O operations. But it would not be possible to hide the fact that the logger library was implemented in terms of the I/O operations, because the handler for the log effects would have to be annotated with I/O operation effects. Furthermore, all log operations would have to be nested in the scope of the log handler, annoyingly inverting control flow relative to the expected approach. And this would have to be done for every library that abstract from the built-in I/O effects, a highly anti-modular approach.

6.1.5 Abstract Algebraic Effects

Our discuss in chapter 4 tackles the issue of abstraction of algebraic effects. This issue was originally raised by Leijen [13], but was not discussed theoretically. Biernacki et al. [2] introduced a core calculus called λ^{HEL} with abstract algebraic effects. However, there are multiple distinctions that distinguish between our core calculus and λ^{HEL} . First, we adopted the agent-based syntax that syntactically distinguishes each module by assigning them to different agents. This design allows us to reason about the code using the information provided by agents, and provide syntactic proof for abstraction properties. Moreover, our calculus can be simply extended with existential types, which serves as an abstraction to represent module systems for a high-level language. The benefit of this design is that the programmer would not need to write embeddings explicitly, as embeddings are generated as a semantic object when a value of existential type is evaluated. In comparison, it is unclear from the paper [2] how a language with a module system could be translated to the coercion-based calculus λ^{HEL} .

Zhang and Myers [32] describe a design for algebraic effects that preserves abstraction in the sense of parametric functions: if a function does not statically know about an algebraic effect, that effect tunnels through that function. This is different from our form of abstraction, in which the definition of an effect is hidden from clients.

Bibliography

- [1] Nada Amin, Tiark Rompf, and Martin Odersky. Foundations of path-dependent types. In *Proceedings of the 2014 ACM International Conference on Object Oriented Programming Systems Languages and Applications*, OOPSLA '14, page 233–249, New York, NY, USA, 2014. Association for Computing Machinery. ISBN 9781450325851. doi: 10.1145/2660193.2660216. URL <https://doi.org/10.1145/2660193.2660216>. 2.1, 2.3.1
- [2] Dariusz Biernacki, Maciej Piròg, Piotr Polesiuk, and Filip Sieczkowski. Abstracting Algebraic Effects. In *Symposium on Principles of Programming Languages*, 2019. 1, 2.2, 4.1, 6.1.5
- [3] Robert L. Bocchino, Jr., Vikram S. Adve, Danny Dig, Sarita V. Adve, Stephen Heumann, Rakesh Komuravelli, Jeffrey Overbey, Patrick Simmons, Hyojin Sung, and Mohsen Vakilian. A Type and Effect System for Deterministic Parallel Java. In *Object Oriented Programming Systems Languages and Applications*, 2009. ISBN 978-1-60558-766-0. doi: 10.1145/1640089.1640097. URL <http://doi.acm.org/10.1145/1640089.1640097>. 1
- [4] JONATHAN IMMANUEL BRACHTHÄUSER, PHILIPP SCHUSTER, and KLAUS OSTERMANN. Effekt: Capability-passing style for type- and effect-safe, extensible effect handlers in scala. *Journal of Functional Programming*, 30:e8, 2020. doi: 10.1017/S0956796820000027. 6.1.4
- [5] Oliver Bračevac, Nada Amin, Guido Salvaneschi, Sebastian Erdweg, Patrick Eugster, and Mira Mezini. Versatile Event Correlation with Algebraic Effects. *Proceedings of the ACM on Programming Languages*, 2(ICFP):67:1–67:31, 2018. ISSN 2475-1421. doi: 10.1145/3236762. URL <http://doi.acm.org/10.1145/3236762>. 1, 2.1, 2.2, 6.1.1
- [6] Stephen Dolan, Spiros Eliopoulos, Daniel Hillerström, Anil Madhavapeddy, K. C. Sivaramakrishnan, and Leo White. Concurrent System Programming with Effect Handlers. In *Trends in Functional Programming*, 2017. doi: 10.1007/978-3-319-89719-6_6. URL https://doi.org/10.1007/978-3-319-89719-6_6. 1
- [7] Andrzej Filinski. Monads in Action. In *Proceedings of the 37th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '10, pages 483–494, New York, NY, USA, 2010. Association for Computing Machinery. ISBN 9781605584799. doi: 10.1145/1706299.1706354. URL <https://doi.org/10.1145/1706299.1706354>. 1, 2.1
- [8] Colin S. Gordon, Werner Dietl, Michael D. Ernst, and Dan Grossman. Javaui: Effects for

- controlling ui object access. In Giuseppe Castagna, editor, *ECOOP 2013 – Object-Oriented Programming*, pages 179–204, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg. ISBN 978-3-642-39038-8. 3.1.1
- [9] Aaron Greenhouse and John Boyland. An object-oriented effects system. pages 205–229, 06 1999. doi: 10.1007/3-540-48743-3_10. 2.1, 3.1.2
- [10] Dan Grossman, Greg Morrisett, and Steve Zdancewic. Syntactic type abstraction. *ACM Trans. Program. Lang. Syst.*, 22(6):1037–1080, November 2000. ISSN 0164-0925. doi: 10.1145/371880.371887. URL <https://doi.org/10.1145/371880.371887>. 1, 4.1, 4.2.2, 5.3.2
- [11] Joseph R. Kiniry. *Advanced Topics in Exception Handling Techniques*, chapter Exceptions in Java and Eiffel: Two Extremes in Exception Design and Application. Springer-Verlag, 2006. ISBN 3-540-37443-4, 978-3-540-37443-5. URL <http://dl.acm.org/citation.cfm?id=2124243.2124264>. 2.1, 6.1.1
- [12] Daan Leijen. Koka: Programming with Row Polymorphic Effect Types. In *Mathematically Structured Functional Programming*, 2014. URL <https://www.microsoft.com/en-us/research/publication/koka-programming-with-row-polymorphic-effect-types-2/>. 2.1, 6.1.3
- [13] Daan Leijen. Algebraic effect handlers with resources and deep finalization. Technical Report MSR-TR-2018-10, April 2018. URL <https://www.microsoft.com/en-us/research/publication/algebraic-effect-handlers-resources-deep-finalization/>. 6.1.5
- [14] K. Rustan M. Leino, Arnd Poetzsch-Heffter, and Yunhong Zhou. Using Data Groups to Specify and Check Side Effects. In *Conference on Programming Language Design and Implementation*, 2002. ISBN 1-58113-463-0. doi: 10.1145/512529.512559. URL <http://doi.acm.org/10.1145/512529.512559>. 2.1
- [15] Paul Levy, John Power, and Hayo Thielecke. Modelling environments in call-by-value programming languages. *Information and Computation*, 185, 08 2003. doi: 10.1016/S0890-5401(03)00088-9. 4.2.1
- [16] Yuheng Long, Yu David Liu, and Hridesh Rajan. Intensional effect polymorphism. In John Tang Boyland, editor, *29th European Conference on Object-Oriented Programming, ECOOP 2015, July 5-10, 2015, Prague, Czech Republic*, volume 37 of *LIPICs*, pages 346–370. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2015. doi: 10.4230/LIPICs.ECOOP.2015.346. URL <https://doi.org/10.4230/LIPICs.ECOOP.2015.346>. 2.1, 6.1.2
- [17] John M. Lucassen. *Types and Effects towards the Integration of Functional and Imperative Programming*. PhD thesis, Massachusetts Institute of Technology, 1987. 2.1, 6.1.1
- [18] John M. Lucassen and David K. Gifford. Polymorphic Effect Systems. In *Symposium on Principles of Programming Languages*, 1988. ISBN 0-89791-252-7. doi: 10.1145/73560.73564. URL <http://doi.acm.org/10.1145/73560.73564>. 1, 2.1, 6.1.1

- [19] Darya Melicher, Yangqingwei Shi, Alex Potanin, and Jonathan Aldrich. A Capability-Based Module System for Authority Control. In *European Conference on Object-Oriented Programming*, 2017. 3.2, 3.2.2, 3.2.6
- [20] John C. Mitchell and Gordon D. Plotkin. Abstract types have existential type. *ACM Trans. Program. Lang. Syst.*, 10(3):470–502, July 1988. doi: 10.1145/44501.45065. URL <https://doi.org/10.1145/44501.45065>. 1
- [21] John C. Mitchell and Gordon D. Plotkin. Abstract types have existential type. *ACM Trans. Program. Lang. Syst.*, 10(3):470–502, July 1988. ISSN 0164-0925. doi: 10.1145/44501.45065. URL <https://doi.org/10.1145/44501.45065>. 5.3
- [22] E. Moggi. Computational lambda-calculus and monads. In *[1989] Proceedings. Fourth Annual Symposium on Logic in Computer Science*, pages 14–23, 1989. doi: 10.1109/LICS.1989.39155. 1
- [23] Martin Odersky and Matthias Zenger. Scalable component abstractions. *SIGPLAN Not.*, 40(10):41–57, October 2005. ISSN 0362-1340. doi: 10.1145/1103845.1094815. URL <https://doi.org/10.1145/1103845.1094815>. 1
- [24] Gordon Plotkin and John Power. Algebraic operations and generic effects. *Applied Categorical Structures*, 11:69–94, 02 2003. doi: 10.1023/A:1023064908962. 4.1
- [25] Gordon Plotkin and John Power. Algebraic Operations and Generic Effects. *Applied Categorical Structures*, 11(1):69–94, 2003. ISSN 1572-9095. doi: 10.1023/A:1023064908962. URL <https://doi.org/10.1023/A:1023064908962>. 2.2
- [26] Gordon Plotkin and Matija Pretnar. Handlers of Algebraic Effects. In *Programming Languages and Systems*, 2009. ISBN 978-3-642-00590-9. 1, 2.2, 4.1
- [27] Lukas Rytz, Martin Odersky, and Philipp Haller. Lightweight Polymorphic Effects. In *European Conference on Object-Oriented Programming*, 2012. ISBN 978-3-642-31056-0. doi: 10.1007/978-3-642-31057-7_13. URL http://dx.doi.org/10.1007/978-3-642-31057-7_13. 2.1, 6.1.3
- [28] Tom Schrijvers, Maciej Piróg, Nicolas Wu, and Mauro Jaskelioff. Monad transformers and modular algebraic effects: What binds them together. In *Proceedings of the 12th ACM SIGPLAN International Symposium on Haskell*, Haskell 2019, page 98–113, New York, NY, USA, 2019. Association for Computing Machinery. ISBN 9781450368131. doi: 10.1145/3331545.3342595. URL <https://doi.org/10.1145/3331545.3342595>. 4.1
- [29] Valery Trifonov and Zhong Shao. Safe and principled language interoperation. In *European Symposium on Programming Languages and Systems*, 1999. 2.1, 6.1.2
- [30] Franklyn A. Turbak and David K. Gifford. *Design Concepts in Programming Languages*. The MIT Press, 2008. ISBN 0262201755, 9780262201759. 1, 2.1
- [31] Marko van Dooren and Eric Steegmans. Combining the robustness of checked exceptions with the flexibility of unchecked exceptions using anchored exception declarations. *SIGPLAN Not.*, 40(10):455–471, October 2005. ISSN 0362-1340. doi: 10.1145/1103845.1094847. URL <https://doi.org/10.1145/1103845.1094847>. 1

- [32] Yizhou Zhang and Andrew C. Myers. Abstraction-safe Effect Handlers via Tunneling. *Proceedings of the ACM on Programming Languages*, 3(POPL):5:1–5:29, 2019. ISSN 2475-1421. doi: 10.1145/3290318. URL <http://doi.acm.org/10.1145/3290318>. 2.2, 6.1.5